

Appendix A

Bibliography

The following references are cited in this book.

- **[A2A Project, A2A 1.0]** A2A Project (Linux Foundation), "Announcing A2A 1.0: JSON-RPC, gRPC, and HTTP+JSON protocol bindings, signed Agent Cards, and multi-tenancy; task-state spelling standardized to 'canceled'," *a2a-protocol.org*, Mar. 2026. [Online]. Available: <https://a2a-protocol.org/latest/announcing-1.0/>
- **[A2A Project, AAIF]** A2A Project, "A new chapter for A2A: joining the Agentic AI Foundation," *A2A Protocol Blog*, Aug. 2026. [Online]. Available: <https://a2a-protocol.org/latest/blog/2026/08/27/a-new-chapter-for-a2a-joining-the-agentic-ai-foundation/>
- **[A2A Project, Specification]** A2A Project (Linux Foundation), "A2A Protocol specification (development version, post-1.0.1; §7.6.4 in-task authorization scope added Jul. 2026)," *a2a-protocol.org*, continuously updated; checked Sep. 2026. [Online]. Available: <https://a2a-protocol.org/dev/specification/>
- **[Advani, Confident Closing]** L. Advani, "From confident closing to silent failure: Characterizing false success in LLM agents," in *FAGEN Workshop at ICML*, Jun. 2026. [Online]. Available: <https://arxiv.org/abs/2606.09863>
- **[AG2, GroupChat]** AG2, "GroupChat: managed multi-agent conversation and speaker-selection patterns," *AG2 Documentation*, continuously updated. [Online]. Available: <https://docs.ag2.ai/latest/docs/user-guide/advanced-concepts/groupchat/groupchat/>
- **[Agent Skills]** Anthropic and community contributors, "Agent Skills: an open standard for portable agent capabilities (SKILL.md)," *agentskills.io*, Dec. 2025. [Online]. Available: <https://agentskills.io/>

- **[Agntcy, Agntcy]** Outshift by Cisco / Linux Foundation, "Agntcy: open infrastructure for the Internet of Agents (OASF, Agent Directory, SLIM, Identity, Observability)," *docs.agntcy.org*, 2025. [Online]. Available: <https://docs.agntcy.org/>
- **[Agrawal et al., GEPA]** L. A. Agrawal, *et al.*, "GEPA: Reflective prompt evolution can outperform reinforcement learning," in *Proc. ICLR Oral*, 2026. [Online]. Available: <https://arxiv.org/abs/2507.19457>
- **[Airbnb, Automation Platform v2]** C. Wang, Z. Xu, P. Lou, *et al.*, "Automation Platform v2: Improving conversational AI at Airbnb," *Airbnb Engineering*, Oct. 2024. [Online]. Available: <https://medium.com/airbnb-engineering/automation-platform-v2-improving-conversational-ai-at-airbnb-d86c9386e0cb>
- **[Airbnb, Eval-Driven Development]** R. Girme, D. Miller, M. Zhao, L. Yang, and C. Kelly, "Eval-driven development: Lessons from evaluating GenAI at scale," *Airbnb Engineering*, Jul. 2026. [Online]. Available: <https://medium.com/airbnb-engineering/eval-driven-development-lessons-from-evaluating-genai-at-scale-e817e5ae5788>
- **[Airbnb, Insight Miner]** W. Dougherty, "Beyond the model: Engineering AI infra with scientific judgement," *Airbnb Engineering*, Sep. 2026. [Online]. Available: <https://medium.com/airbnb-engineering/beyond-the-model-engineering-ai-infra-with-scientific-judgement-371316d43261>
- **[Airbnb, Intelligent Automation Platform]** Z. Xu, *et al.*, "Intelligent Automation Platform: Empowering conversational AI and beyond at Airbnb," *Airbnb Engineering*, Jan. 2022. [Online]. Available: <https://medium.com/airbnb-engineering/intelligent-automation-platform-empowering-conversational-ai-and-beyond-at-airbnb-869c44833ff2>
- **[Airbnb, Q2 2025 Earnings Call]** Airbnb, Inc., "Q2 2025 earnings call transcript" (corrected transcript, FactSet CallStreet), *Airbnb Investor Relations*, Aug. 2025. [Online]. Available: https://s26.q4cdn.com/656283129/files/doc_financials/2025/q2/Airbnb-Q2-25-Earnings-Call-Transcript.pdf
- **[Airbnb, Q3 2025 Letter]** Airbnb, Inc., "Q3 2025 shareholder letter," *Airbnb Investor Relations*, Nov. 2025. [Online]. Available: https://s26.q4cdn.com/656283129/files/doc_financials/2025/q3/v2/Airbnb_Q3-2025-Shareholder-Letter.pdf
- **[Airbnb, Shareholder Letter]** Airbnb, Inc., "Q2 2026 shareholder letter," *Airbnb Investor Relations*, Aug. 2026. [Online]. Available: https://s26.q4cdn.com/656283129/files/doc_financials/2026/q2/Airbnb-Q2-2026-Shareholder-Letter.pdf

- **[Alshahwan et al., TestGen-LLM]** N. Alshahwan, J. Chheda, A. Finogenova, *et al.*, "Automated unit test improvement using large language models at Meta," in *Companion Proc. 32nd ACM Int. Conf. Foundations of Software Engineering (FSE)*, Jul. 2024, pp. 185–196. [Online]. Available: <https://doi.org/10.1145/3663529.3663839>
- **[Altmann et al., Emergence in multi-agent]** P. Altmann, *et al.*, "Emergence in multi-agent systems: A safety perspective," *arXiv:2408.04514*, Aug. 2024. [Online]. Available: <https://arxiv.org/abs/2408.04514>
- **[Andon Labs, Vending-Bench Arena]** Andon Labs, "Opus 5 on Vending-Bench: Once again the best capitalist, once again misaligned," *Andon Labs Blog*, Jul. 2026. [Online]. Available: <https://andonlabs.com/blog/opus-5-vending-bench>
- **[Andriushchenko et al., AgentHarm]** M. Andriushchenko, *et al.*, "AgentHarm: A benchmark for measuring harmfulness of LLM agents," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.09024>
- **[ANP Community, DID-based agent discovery]** ANP Working Group, "Agent Network Protocol specification (DID-based agent discovery)," 2024–. [Online]. Available: <https://github.com/agent-network-protocol/AgentNetworkProtocol>
- **[Anthropic, Advanced Tool Use]** Anthropic, "Introducing advanced tool use on the Claude Developer Platform," *Anthropic Engineering*, Nov. 2025. [Online]. Available: <https://www.anthropic.com/engineering/advanced-tool-use>
- **[Anthropic, Agent Autonomy]** M. McCain, T. Millar, S. Huang, *et al.* (Anthropic), "Measuring AI agent autonomy in practice," *Anthropic Research*, Feb. 2026. [Online]. Available: <https://www.anthropic.com/research/measuring-agent-autonomy>
- **[Anthropic, AI-Native Stacks Panel]** B. Robertson (Anthropic) with N. Grouppen (Harvey), W. Yan (Cognition), and D. Fatihah (Gamma), "Building AI-native: Inside the stacks powering Cognition, Gamma, and Harvey," *Code with Claude 2026 (YouTube)*, May 2026. [Online]. Available: <https://www.youtube.com/watch?v=0FDm3T7pV1c>
- **[Anthropic, ASL-3]** Anthropic, "AI Safety Level 3 deployment safeguards report," *Anthropic*, May 2025. [Online]. Available: <https://www.anthropic.com/asl3-deployment-safeguards>
- **[Anthropic, Auto Mode]** J. Hughes (Anthropic), "How we built Claude Code auto mode: A safer way to skip permissions," *Anthropic Engineering*, Mar. 2026. [Online]. Available: <https://www.anthropic.com/engineering/claude-code-auto-mode>

- **[Anthropic, Batch Processing]** Anthropic, "Batch processing — the Message Batches API," *Anthropic Platform Documentation*, continuously updated. [Online]. Available: <https://platform.claude.com/docs/en/build-with-claude/batch-processing>
- **[Anthropic, Building Effective Agents]** E. Schluntz and B. Zhang (Anthropic), "Building effective agents," *Anthropic Research*, Dec. 2024. [Online]. Available: <https://www.anthropic.com/research/building-effective-agents>
- **[Anthropic, Built-in Memory]** Anthropic, "Built-in memory for Claude Managed Agents," *Claude Blog*, Apr. 2026. [Online]. Available: <https://claude.com/blog/claude-managed-agents-memory>
- **[Anthropic, C Compiler]** N. Carlini (Anthropic), "Building a C compiler with a team of parallel Claudes," *Anthropic Engineering*, Feb. 2026. [Online]. Available: <https://www.anthropic.com/engineering/building-c-compiler>
- **[Anthropic, Classifiers++]** H. Cunningham, J. Wei, *et al.* (Anthropic), "Constitutional Classifiers++: Efficient production-grade defenses against universal jailbreaks," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2601.04603>
- **[Anthropic, Claude Agent SDK]** Anthropic, "Claude Agent SDK: agentic coding + MCP tools," *Anthropic Platform Docs*, 2025. [Online]. Available: <https://platform.claude.com/docs/en/agent-sdk/overview>
- **[Anthropic, Claude Code Checkpointing]** Anthropic, "Checkpointing: track, rewind, and summarize Claude's edits and conversation to manage session state," *Claude Code Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://code.claude.com/docs/en/checkpointing>
- **[Anthropic, Claude Code Context Window]** Anthropic, "Explore the context window: what loads at startup and what survives compaction," *Claude Code Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://code.claude.com/docs/en/context-window>
- **[Anthropic, Claude Code Costs]** Anthropic, "Manage costs effectively: per-developer averages, cost drivers, and levers for Claude Code," *Claude Code Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://code.claude.com/docs/en/costs>
- **[Anthropic, Claude Code Hooks]** Anthropic, "Hooks reference: hook events, exit codes, and decision control," *Claude Code Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://code.claude.com/docs/en/hooks>

- **[Anthropic, Claude Code Memory]** Anthropic, "How Claude remembers your project: CLAUDE.md, AGENTS.md, and auto memory," *Claude Code Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://code.claude.com/docs/en/memory>
- **[Anthropic, Claude Code Monitoring]** Anthropic, "Monitoring: OpenTelemetry metrics and events for Claude Code," *Claude Code Documentation*, continuously updated. [Online]. Available: <https://code.claude.com/docs/en/monitoring-usage>
- **[Anthropic, Claude Code Permission Modes]** Anthropic, "Choose a permission mode: auto mode, its classifier, and how it handles subagents," *Claude Code Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://code.claude.com/docs/en/permission-modes>
- **[Anthropic, Claude Code Permissions]** Anthropic, "Configure permissions: rule evaluation order, sandboxing, and workspace trust," *Claude Code Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://code.claude.com/docs/en/permissions>
- **[Anthropic, Claude Code Subagents]** Anthropic, "Create custom subagents," *Claude Code Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://code.claude.com/docs/en/sub-agents>
- **[Anthropic, Claude Code Workflows]** Anthropic, "Orchestrate subagents at scale with dynamic workflows," *Claude Code Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://code.claude.com/docs/en/workflows>
- **[Anthropic, Code Execution MCP]** Anthropic, "Code execution with MCP: Building more efficient agents," *Anthropic Engineering*, Nov. 2025. [Online]. Available: <https://www.anthropic.com/engineering/code-execution-with-mcp>
- **[Anthropic, Coding Trends]** Anthropic, "2026 Agentic Coding Trends Report," *Anthropic*, 2026. [Online]. Available: <https://resources.anthropic.com/hubfs/2026%20Agentic%20Coding%20Trends%20Report.pdf>
- **[Anthropic, Containing Claude]** Anthropic, "How we contain Claude across products," *Anthropic Engineering*, May 2026. [Online]. Available: <https://www.anthropic.com/engineering/how-we-contain-claude>
- **[Anthropic, Context Engineering]** Anthropic Applied AI Team, "Effective context engineering for AI agents," *Anthropic Engineering*, Sep. 2025. [Online]. Available: <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>

- **[Anthropic, Context Management]** Anthropic, "Managing context on the Claude Developer Platform," *Anthropic News*, Sep. 2025. [Online]. Available: <https://claude.com/blog/context-management>
- **[Anthropic, Cybersecurity Incidents]** Anthropic, "An alignment assessment of recent cybersecurity incidents," *Anthropic Research*, Sep. 2026. [Online]. Available: <https://www.anthropic.com/research/alignment-assessment-cybersecurity-incidents>
- **[Anthropic, Demystifying Evals]** Anthropic, "Demystifying evals for AI agents," *Anthropic Engineering*, Jan. 2026. [Online]. Available: <https://www.anthropic.com/engineering/demystifying-evals-for-ai-agents>
- **[Anthropic, Effective Tools]** Anthropic, "Writing effective tools for agents — with agents," *Anthropic Engineering*, Sep. 2025. [Online]. Available: <https://www.anthropic.com/engineering/writing-tools-for-agents>
- **[Anthropic, Extended Thinking]** Anthropic, "Thinking: extended thinking and redacted thinking blocks," *Anthropic API Documentation*, continuously updated. [Online]. Available: <https://platform.claude.com/docs/en/build-with-claude/thinking#redacted-thinking-blocks>
- **[Anthropic, Infrastructure Noise]** G. Segato (Anthropic), "Quantifying infrastructure noise in agentic coding evals," *Anthropic Engineering*, Feb. 2026. [Online]. Available: <https://www.anthropic.com/engineering/infrastructure-noise>
- **[Anthropic, MCP]** D. Soria Parra and J. Spahr-Summers, "Introducing the Model Context Protocol," *Anthropic*, Nov. 2024. [Online]. Available: <https://www.anthropic.com/news/model-context-protocol>
- **[Anthropic, Memory and Dreaming]** Anthropic (M. Murag, Platform Team), "Memory and dreaming for self-learning agents," *YouTube (Code with Claude 2026)*, May 2026. [Online]. Available: <https://www.youtube.com/watch?v=RtywqDFBYnQ>
- **[Anthropic, Multi-Agent Research]** J. Hadfield, B. Zhang, K. Lien, B. Scholz, B. Fox, and D. Ford (Anthropic), "How we built our multi-agent research system," *Anthropic Engineering*, Jun. 2025. [Online]. Available: <https://www.anthropic.com/engineering/multi-agent-research-system>
- **[Anthropic, Poisoning]** Anthropic, UK AI Security Institute, and Alan Turing Institute, "A small number of samples can poison LLMs of any size," *Anthropic Research*, Oct. 2025. [Online]. Available: <https://www.anthropic.com/research/small-samples-poison>

- **[Anthropic, Postmortem]** Anthropic, "A postmortem of three recent issues," *Anthropic Engineering*, Sep. 2025. [Online]. Available: <https://www.anthropic.com/engineering/a-postmortem-of-three-recent-issues>
- **[Anthropic, Pricing]** Anthropic, "Pricing: model rates, prompt-caching multipliers, tool-use system-prompt tokens, and server-side tool fees," *Claude API Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://platform.claude.com/docs/en/about-claude/pricing>
- **[Anthropic, Prompt Caching]** Anthropic, "Prompt caching with Claude," *Anthropic News*, Aug. 2024. [Online]. Available: <https://claude.com/blog/prompt-caching>
- **[Anthropic, Prompt Caching Docs]** Anthropic, "Prompt caching: usage fields, cache breakpoints, and invalidation," *Claude API Documentation*, continuously updated. [Online]. Available: <https://platform.claude.com/docs/en/build-with-claude/prompt-caching>
- **[Anthropic, Prompt Injection Defenses]** Anthropic, "Mitigating the risk of prompt injections in browser use," *Anthropic Research*, Nov. 2025. [Online]. Available: <https://www.anthropic.com/research/prompt-injection-defenses>
- **[Anthropic, Usage and Cost API]** Anthropic, "Usage and Cost API," *Claude API Documentation*, continuously updated. [Online]. Available: <https://platform.claude.com/docs/en/manage-claude/usage-cost-api>
- **[AP2, Specification]** Agent Payments Protocol authors, "Agentic Payment Protocol v0.2 specification: checkout and payment mandates, autonomous constraints, receipts, and dispute evidence," *AP2 Documentation*, 2026. [Online]. Available: <https://github.com/google-agentic-commerce/AP2/blob/main/docs/ap2/specification.md>
- **[Aradhyula et al., Unlocking 10× productivity]** S. Aradhyula, *et al.*, "Unlocking 10× productivity: A blueprint for distributed multi-agentic AI platform engineering," *Outshift by Cisco*, Dec. 2025. [Online]. Available: https://outshift-headless-cms-s3.s3.us-east-2.amazonaws.com/CAIPE_WP_f813d70ad8.pdf
- **[Arcuschin et al., Unfaithful CoT]** I. Arcuschin, J. Janiak, R. Krzyzanowski, S. Rajamanoharan, N. Nanda, and A. Conmy, "Chain-of-thought reasoning in the wild is not always faithful," in *Proc. ICML*, 2026. [Online]. Available: <https://arxiv.org/abs/2503.08679>
- **[Asawa et al., Continual Learning Bench]** P. Asawa, C. M. Glaze, G. Orlanski, *et al.*, "Continual Learning Bench: Evaluating frontier AI systems in real-world stateful

- environments," in *Proc. NeurIPS Evaluations and Datasets Track*, 2026. [Online]. Available: <https://arxiv.org/abs/2606.05661>
- **[ASB, ASB]** H. Zhang, *et al.*, "Agent Security Bench (ASB): Formalizing and benchmarking attacks and defenses in LLM-based agents," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.02644>
 - **[Ashery et al., Social conventions]** A. F. Ashery, L. M. Aiello, and A. Baronchelli, "Emergent social conventions and collective bias in LLM populations," *Science Advances*, vol. 11, no. 20, eadu9368, May 2025. [Online]. Available: <https://www.science.org/doi/10.1126/sciadv.adu9368>
 - **[Atil et al., Non-determinism of 'deterministic']** B. Atil, *et al.*, "Non-determinism of 'deterministic' LLM settings," *arXiv:2408.04667*, Aug. 2024. [Online]. Available: <https://arxiv.org/abs/2408.04667>
 - **[Atla, Selene Mini]** Atla, "Atla Selene Mini (Selene-1-Mini-Llama-3.1-8B): a general-purpose small evaluation model," *Hugging Face model card*, Jan. 2025. [Online]. Available: <https://huggingface.co/AtlaAI/Selene-1-Mini-Llama-3.1-8B>
 - **[Atlan, AI agent harness]** Atlan, "AI agent harness failures: 13 anti-patterns and root causes," *Atlan Industry Analysis*, 2026. [Online]. Available: <https://atlan.com/know/agent-harness-failures-anti-patterns/>
 - **[AWS Bedrock Agents]** Amazon Web Services, "Amazon Bedrock AgentCore — the platform for production AI agents, any framework, any model (successor to Amazon Bedrock Agents, now Bedrock Agents Classic and closed to new customers)," *AWS Product Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://aws.amazon.com/bedrock/agentcore/>
 - **[AWS DynamoDB, Amazon DynamoDB]** Amazon Web Services, "Amazon DynamoDB — eventually consistent key-value store," *AWS Product Documentation*, continuously updated. [Online]. Available: <https://aws.amazon.com/dynamodb/>
 - **[AWS Lambda]** Amazon Web Services, "AWS Lambda — serverless compute service," *AWS Product Documentation*, continuously updated. [Online]. Available: <https://docs.aws.amazon.com/lambda/>
 - **[AWS Secrets Manager, AWS Secrets Manager]** Amazon Web Services, "AWS Secrets Manager — managed rotation, retrieval, and lifecycle for secrets," *AWS Product Documentation*, continuously updated. [Online]. Available: <https://docs.aws.amazon.com/secretsmanager/>

- **[AWS SQS, SQS]** Amazon Web Services, "Amazon Simple Queue Service (SQS) — managed message-queueing for distributed systems," *AWS Product Documentation*, continuously updated. [Online]. Available: <https://docs.aws.amazon.com/AWSSimpleQueueService/>
- **[AWS Step Functions, AWS Step Functions]** Amazon Web Services, "AWS Step Functions — durable workflow orchestration as code," *AWS Product Documentation*, continuously updated. [Online]. Available: <https://docs.aws.amazon.com/step-functions/>
- **[AWS, AgentCore Policy]** L. Hadarean and J.-B. Tristan (AWS), "Why Policy in Amazon Bedrock AgentCore chose Cedar for securing agentic workflows," *AWS Security Blog*, May 2026. [Online]. Available: <https://aws.amazon.com/blogs/security/why-policy-in-amazon-bedrock-agentcore-chose-cedar-for-securing-agentic-workflows/>
- **[AWS, Bedrock Guardrails]** Amazon Web Services, "Amazon Bedrock Guardrails: contextual grounding check and Automated Reasoning checks," *Amazon Bedrock User Guide*, continuously updated. [Online]. Available: <https://docs.aws.amazon.com/bedrock/latest/userguide/guardrails.html>
- **[AWS, Durable Functions]** R. Pisal and B. Freiberg (Amazon Web Services), "Building fault-tolerant applications with AWS Lambda durable functions," *AWS Compute Blog*, Feb. 2026. [Online]. Available: <https://aws.amazon.com/blogs/compute/building-fault-tolerant-long-running-application-with-aws-lambda-durable-functions/>
- **[AWS, Strands]** AWS, "Customize agent workflows with advanced orchestration techniques using Strands Agents," *AWS Machine Learning Blog*, 2025. [Online]. Available: <https://aws.amazon.com/blogs/machine-learning/customize-agent-workflows-with-advanced-orchestration-techniques-using-strands-agents/>
- **[Azure AI Agent Service, Azure AI Agent]** Microsoft, "Microsoft Foundry Agent Service — managed agent runtime on Microsoft Foundry (formerly Azure AI Foundry Agent Service, and before that Azure AI Agent Service)," *Microsoft Learn*, continuously updated. [Online]. Available: <https://learn.microsoft.com/en-us/azure/foundry/agents/overview>
- **[Baddeley & Hitch, Working memory]** A. D. Baddeley and G. Hitch, "Working memory," in *The Psychology of Learning and Motivation*, G. H. Bower, Ed. New York: Academic Press, 1974, vol. 8, pp. 47–89.

- **[Bagdasarian et al., AirGapAgent]** E. Bagdasarian, R. Yi, S. Ghalebikesabi, P. Kairouz, M. Gruteser, S. Oh, B. Balle, and D. Ramage, "AirGapAgent: Protecting privacy-conscious conversational agents," in *Proc. ACM CCS*, 2024. [Online]. Available: <https://arxiv.org/abs/2405.05175>
- **[Bai et al., Constitutional AI]** Y. Bai, S. Kadavath, S. Kundu, *et al.*, "Constitutional AI: Harmlessness from AI feedback," *arXiv:2212.08073*, Dec. 2022. [Online]. Available: <https://arxiv.org/abs/2212.08073>
- **[Bai et al., FVDebug]** Y. Bai, G. Bany Hamad, C.-T. Ho, S. Suhaib, and H. Ren, "FVDebug: An LLM-driven debugging assistant for automated root cause analysis of formal verification failures," in *Proc. DVCon U.S.*, 2026. [Online]. Available: <https://arxiv.org/abs/2510.15906>
- **[Banik et al., All Smoke]** D. Banik, K. Chowdhury, and S. I. Shamim, "All smoke, no alarm: Oracle signals in agent-authored test code," in *Proc. 8th IEEE Int. Conf. Artificial Intelligence Testing (AITest)*, short paper, Jul. 2026. [Online]. Available: <https://arxiv.org/abs/2606.18168>
- **[Bansal et al., accuracy]** G. Bansal, *et al.*, "Beyond accuracy: The role of mental models in human-AI team performance," in *Proc. HCOMP*, 2019. [Online]. Available: <https://ojs.aaai.org/index.php/HCOMP/article/view/5285>
- **[Bansal et al., Magentic Marketplace]** G. Bansal, W. Hua, Z. Huang, *et al.* (Microsoft Research), "Magentic Marketplace: An open-source environment for studying agentic markets," *Microsoft Research tech. rep. MSR-TR-2025-50*, Oct. 2025. [Online]. Available: <https://arxiv.org/abs/2510.25779>
- **[Barres et al., T²-bench]** V. Barres, *et al.*, "T²-bench: Evaluating conversational agents in a dual-control environment," in *Proc. ICML*, 2026. [Online]. Available: <https://arxiv.org/abs/2506.07982>
- **[Batra et al., SALT]** S. Batra, P. Tillman, S. Gaggar, *et al.*, "SALT: Steering activations towards leakage-free thinking in chain of thought," in *Proc. NeurIPS Workshop on Socially Responsible and Trustworthy Foundation Models*, 2025. [Online]. Available: <https://arxiv.org/abs/2511.07772>
- **[Becker et al., Early-2025 AI]** J. Becker, N. Rush, E. Barnes, and D. Rein, "Measuring the impact of early-2025 AI on experienced open-source developer productivity," *arXiv: 2507.09089*, Jul. 2025. [Online]. Available: <https://arxiv.org/abs/2507.09089>

- **[Berner et al., OpenAI Five]** C. Berner, G. Brockman, B. Chan, *et al.*, "Dota 2 with large scale deep reinforcement learning," *arXiv:1912.06680*, Dec. 2019. [Online]. Available: <https://arxiv.org/abs/1912.06680>
- **[Beurer-Kellner et al., Design Patterns]** L. Beurer-Kellner, B. Buesser, A.-M. Crețu, E. Debenedetti, D. Dobos, D. Fabian, M. Fischer, D. Froelicher, K. Grosse, D. Naeff, E. Ozoani, A. Paverd, F. Tramèr, and V. Volhejn, "Design patterns for securing LLM agents against prompt injections," *arXiv:2506.08837*, Jun. 2025. [Online]. Available: <https://arxiv.org/abs/2506.08837>
- **[Beyer et al., Site Reliability Engineering]** B. Beyer, C. Jones, J. Petoff, and N. R. Murphy, *Site Reliability Engineering: How Google Runs Production Systems*. Sebastopol, CA: O'Reilly Media, 2016. [Online]. Available: <https://sre.google/sre-book/table-of-contents/>
- **[Bjork, Desirable Difficulty]** R. A. Bjork, "Memory and metamemory considerations in the training of human beings," in *Metacognition: Knowing About Knowing*, J. Metcalfe and A. P. Shimamura, Eds. Cambridge, MA: MIT Press, 1994, pp. 185–205.
- **[Borrett et al., Competing with AI Scientists]** T. Borrett, L. Xu, A. Nilipour, B. Bolliet, *et al.*, "Competing with AI scientists: Agent-driven approach to astrophysics research," *arXiv:2604.09621*, Mar. 2026. [Online]. Available: <https://arxiv.org/abs/2604.09621>
- **[Bragg et al., AstaBench]** J. Bragg, *et al.*, "AstaBench: Rigorous benchmarking of AI agents with a scientific research suite," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2510.21652>
- **[Branwen, Scaling Hypothesis]** G. Branwen, "The scaling hypothesis," *gwern.net*, 2020. [Online]. Available: <https://gwern.net/scaling-hypothesis>
- **[Breunig, long contexts fail]** D. Breunig, "How long contexts fail (poison / distraction / confusion / clash taxonomy)," *dbreunig.com*, Jun. 2025. [Online]. Available: <https://www.dbreunig.com/2025/06/22/how-contexts-fail-and-how-to-fix-them.html>
- **[Brooker, Exponential backoff]** M. Brooker, "Exponential backoff and jitter," *AWS Architecture Blog*, Mar. 2015. [Online]. Available: <https://aws.amazon.com/blogs/architecture/exponential-backoff-and-jitter/>
- **[Brooks, Mythical Man-Month]** F. P. Brooks, *The Mythical Man-Month: Essays on Software Engineering*. Reading, MA: Addison-Wesley, 1975. [Online]. Available: https://en.wikipedia.org/wiki/The_Mythical_Man-Month

- **[Brown et al., Large Language Monkeys]** B. Brown, J. Juravsky, R. Ehrlich, *et al.*, "Large language monkeys: Scaling inference compute with repeated sampling," *arXiv*: 2407.21787, Dec. 2024. [Online]. Available: <https://arxiv.org/abs/2407.21787>
- **[Brown, Agent Swarms]** N. Brown (OpenAI), interviewed by D. Patel, "OpenAI researcher on agent swarms & recursive self-improvement," *Dwarkesh Podcast (YouTube)*, Sep. 2026. [Online]. Available: <https://www.youtube.com/watch?v=6AgOfiZOWiY>
- **[CAIPE, Agent Evaluation]** S. Aradhyula, "Initiative: Agent & RAG evaluation," *CAIPE GitHub Issues*, no. 2081, Jun. 2026. [Online]. Available: <https://github.com/caipe-io/ai-platform-engineering/issues/2081>
- **[CAIPE, Agent Memory]** S. Aradhyula, "Initiative: Agent memory & durable user context," *CAIPE GitHub Issues*, no. 2486, Aug. 2026. [Online]. Available: <https://github.com/caipe-io/ai-platform-engineering/issues/2486>
- **[CAIPE, Docs]** CAIPE project, "Solution architecture," *caipe.io documentation*, release 1.0.0, Sep. 2026. [Online]. Available: <https://caipe.io/docs/1.0.0/architecture/>
- **[CAIPE, RBAC]** CAIPE project, "Security: role-based access control," *caipe.io documentation*, release 1.0.0, checked Sep. 2026. [Online]. Available: <https://caipe.io/docs/1.0.0/security/rbac/>
- **[CAIPE, Release 1.1.0]** CAIPE project, "Release 1.1.0: Automation, access, and runtime updates," *caipe.io release notes*, Sep. 2026. [Online]. Available: <https://caipe.io/releases/release-1.1.0>
- **[CAIPE, Release 1.1.1]** CAIPE project, "Release 1.1.1: Helm and service account fixes," *caipe.io release notes*, Sep. 2026. [Online]. Available: <https://caipe.io/releases/release-1.1.1>
- **[CAIPE, Remote Agent Tool]** S. Aradhyula, "feat(dynamic-agents): integrate remote agents as tool calls via A2A protocol," *CAIPE GitHub Issues*, no. 2013, Jun. 2026. [Online]. Available: <https://github.com/caipe-io/ai-platform-engineering/issues/2013>
- **[CAIPE, Source]** CAIPE project, "caipe-io/ai-platform-engineering: source code and in-repository documentation," *GitHub*, release 1.0.0 (Sep. 2026), read with its history back to release 0.5.21 and forward to release 1.1.1. [Online]. Available: <https://github.com/caipe-io/ai-platform-engineering/tree/1.0.0>
- **[CAIPE, Supervisor Deprecation]** S. Aradhyula, "Deprecating A2A supervisor and packaged A2A sub-agents," *CAIPE GitHub Discussions*, no. 1997, Jun. 2026. [Online]. Available: <https://github.com/orgs/caipe-io/discussions/1997>

- **[Carlini et al., AutoAdvExBench]** N. Carlini, *et al.*, "AutoAdvExBench: Benchmarking autonomous exploitation of adversarial-example defenses," in *Proc. ICML, 2025 (ICML 2025)*. [Online]. Available: <https://arxiv.org/abs/2503.01811>
- **[Carlini et al., Memorization]** N. Carlini, D. Ippolito, M. Jagielski, K. Lee, F. Tramèr, and C. Zhang, "Quantifying memorization across neural language models," in *Proc. ICLR, 2022*. [Online]. Available: <https://arxiv.org/abs/2202.07646>
- **[CCPA, Right to Delete]** California Legislature, "California Consumer Privacy Act, Cal. Civ. Code § 1798.105: Consumers' Right to Delete Personal Information," *California Legislative Information*. [Online]. Available: https://leginfo.legislature.ca.gov/faces/codes_displaySection.xhtml?lawCode=CIV§ionNum=1798.105
- **[Cemri et al., MAST]** M. Cemri, M. Z. Pan, S. Yang, *et al.*, "Why do multi-agent LLM systems fail? (MAST)," in *Proc. NeurIPS Datasets and Benchmarks Track, 2025*. [Online]. Available: <https://arxiv.org/abs/2503.13657>
- **[Chan et al., ChatEval]** C.-M. Chan, *et al.*, "ChatEval: Towards better LLM-based evaluators through multi-agent debate," in *Proc. ICLR, 2024*. [Online]. Available: <https://arxiv.org/abs/2308.07201>
- **[Chan et al., MLE-bench]** J. S. Chan, N. Chowdhury, O. Jaffe, *et al.*, "MLE-bench: Evaluating machine learning agents on machine learning engineering," in *Proc. ICLR, 2025*. [Online]. Available: <https://arxiv.org/abs/2410.07095>
- **[Chang et al., DRC-Coder]** C.-C. Chang, C.-T. Ho, Y. Li, Y. Chen, and H. Ren, "DRC-Coder: Automated DRC checker code generation using LLM autonomous agent," in *Proc. ACM Int. Symp. Physical Design (ISPD)*, Mar. 2025, pp. 143–151. [Online]. Available: <https://arxiv.org/abs/2412.05311>
- **[Chaudhry et al., Murakkab]** G. I. Chaudhry, E. Choukse, H. Qiu, Í. Goiri, R. Fonseca, A. Belay, and R. Bianchini, "Murakkab: Resource-efficient agentic workflow orchestration in cloud platforms," in *Proc. OSDI, 2026*, pp. 567–587. [Online]. Available: <https://www.usenix.org/conference/osdi26/presentation/chaudhry>
- **[Chen et al., AgentSlimming]** Y. Chen, *et al.*, "AgentSlimming: Towards efficient and cost-aware multi-agent systems," in *Proc. ACL, 2026*. [Online]. Available: <https://arxiv.org/abs/2605.08813>
- **[Chen et al., Codex]** M. Chen, J. Tworek, H. Jun, *et al.*, "Evaluating large language models trained on code (Codex / HumanEval)," *arXiv:2107.03374*, Jul. 2021. [Online]. Available: <https://arxiv.org/abs/2107.03374>

- **[Chen et al., Overthinking]** X. Chen, J. Xu, T. Liang, *et al.*, "Do NOT think that much for $2+3=?$ On the overthinking of long reasoning models," in *Proc. ICML*, 2025. [Online]. Available: <https://arxiv.org/abs/2412.21187>
- **[Chen et al., Whole Elephant]** M. Chen, J. Wang, F. Mu, Y. Wang, Z. Liu, H. Feng, and Q. Wang, "Seeing the whole elephant: A benchmark for failure attribution in LLM-based multi-agent systems," in *Proc. ACL*, 2026. [Online]. Available: <https://arxiv.org/abs/2604.22708>
- **[Cheng et al., Self-Adaptive Roadmap]** B. H. C. Cheng, R. de Lemos, H. Giese, *et al.*, "Software engineering for self-adaptive systems: A research roadmap," in *Software Engineering for Self-Adaptive Systems, LNCS 5525*, Springer, 2009, pp. 1–26. [Online]. Available: https://doi.org/10.1007/978-3-642-02161-9_1
- **[Chiang et al., Chatbot Arena]** W.-L. Chiang, *et al.*, "Chatbot Arena: An open platform for evaluating LLMs by human preference," in *Proc. ICML*, 2024. [Online]. Available: <https://arxiv.org/abs/2403.04132>
- **[Choi et al., Debate or Vote]** H. K. Choi, X. Zhu, and S. Li, "Debate or vote: Which yields better decisions in multi-agent large language models?" in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2508.17536>
- **[Chow et al., PhysBench]** W. Chow, J. Mao, B. Li, D. Seita, V. C. Guizilini, and Y. Wang, "PhysBench: Benchmarking and enhancing vision-language models for physical world understanding," in *Proc. ICLR*, 2025. [Online]. Available: <https://openreview.net/forum?id=Q6a9W6kzv5>
- **[Chowdhury et al., Code Review Agents]** K. Chowdhury, D. Banik, K. M. Ferdous, and S. I. Shamim, "From industry claims to empirical reality: An empirical study of code review agents in pull requests," in *Proc. 23rd Int. Conf. Mining Software Repositories (MSR)*, Apr. 2026. [Online]. Available: <https://arxiv.org/abs/2604.03196>
- **[Christiano et al., RLHF]** P. F. Christiano, J. Leike, T. Brown, M. Martic, S. Legg, and D. Amodei, "Deep reinforcement learning from human preferences," in *Proc. NeurIPS*, 2017. [Online]. Available: <https://arxiv.org/abs/1706.03741>
- **[Cisco, AI PODs]** Cisco Systems, "Cisco AI PODs: pre-validated, modular infrastructure for enterprise AI, combining Cisco UCS compute, Nexus networking, NVIDIA GPUs, and NVIDIA AI Enterprise in training- and inference-sized configurations," *Cisco solution page*, continuously updated; checked Sep. 2026. [Online]. Available: <https://www.cisco.com/site/us/en/solutions/artificial-intelligence/infrastructure/ai-pods.html>

- **[Cisco, Cisco AI Defense]** Cisco Systems, "Cisco AI Defense (agentic expansion: AI BOM, MCP Catalog, A2A Scanner, adaptive algorithmic red-teaming, NeMo Guardrails integration); DefenseClaw open-source agentic-runtime security framework," *Cisco product page*, RSAC 2026. [Online]. Available: <https://www.cisco.com/site/us/en/products/security/ai-defense/index.html>
- **[Clark et al., SREGym]** J. Clark, Y. Su, *et al.*, "SREGym: A live benchmark for AI SRE agents with high-fidelity failure scenarios," *arXiv:2605.07161*, May 2026. [Online]. Available: <https://arxiv.org/abs/2605.07161>
- **[Claude Cowork incident 2026]** N. Khandusenko, "A man asked Claude Cowork to tidy up his wife's files. A few minutes later, he almost had a heart attack," *dev.ua*, 9 Feb. 2026, reporting Nick Davidov's own account (spelled Davydov by dev.ua) of the 7 Feb. 2026 incident: permission granted only for temporary Office files, a rename that deleted the photo folder, "there are no files in the trash because they were deleted through the terminal," and "tens of thousands of files" restored from iCloud Drive's 30-day window. Corroborated by *Futurism*, 13 Feb. 2026, and *Dexerto*, 16 Feb. 2026. Anthropic published no postmortem for this event. [Online]. Available: <https://dev.ua/en/news/claude-cowork-1770633970>
- **[CMBAgent, cmbagent]** CMBAgent Contributors, "cmbagent: Open-source multi-agent system for cosmology science," *GitHub*, 2026. [Online]. Available: <https://github.com/CMBAgents/cmbagent>
- **[Cobbe et al., Training Verifiers]** K. Cobbe, V. Kosaraju, M. Bavarian, *et al.*, "Training verifiers to solve math word problems," *arXiv:2110.14168*, Nov. 2021. [Online]. Available: <https://arxiv.org/abs/2110.14168>
- **[Cognition, Don't build multi-agents]** W. Yan (Cognition), "Don't build multi-agents," *cognition.ai*, Jun. 2025. [Online]. Available: <https://cognition.com/blog/dont-build-multi-agents>
- **[Cognition, Multi-agents]** W. Yan (Cognition), "Multi-agents: What's actually working," *cognition.ai*, Apr. 2026. [Online]. Available: <https://cognition.com/blog/multi-agents-working>
- **[Crandall et al., Working Minds]** B. Crandall, G. Klein, and R. R. Hoffman, *Working Minds: A Practitioner's Guide to Cognitive Task Analysis*. Cambridge, MA: MIT Press, 2006. [Online]. Available: <https://mitpress.mit.edu/9780262532815/working-minds/>

- **[CrewAI, Cognitive Memory]** J. Moura (CrewAI), "How we built Cognitive Memory for Agentic Systems," *CrewAI Blog*, Mar. 2026. [Online]. Available: <https://www.crewai.com/blog/how-we-built-cognitive-memory-for-agentic-systems>
- **[crewAI, CrewAI]** J. Moura, "CrewAI: Role-based agent orchestration framework," *GitHub*, 2024. [Online]. Available: <https://github.com/crewAIInc/crewAI>
- **[CSA, AAGATE]** Cloud Security Alliance, "AAGATE (Agentic AI Governance Assurance and Trust Engine): A NIST AI RMF-aligned governance platform for agentic AI," *CSA Blog*, Dec. 2025. [Online]. Available: <https://cloudsecurityalliance.org/blog/2025/12/22/aagate-a-nist-ai-rmf-aligned-governance-platform-for-agentic-ai>
- **[CSA, MAESTRO]** Cloud Security Alliance, "Agentic AI threat modeling framework: MAESTRO (seven-layer reference model)," *CSA Blog*, Feb. 2025. [Online]. Available: <https://cloudsecurityalliance.org/blog/2025/02/06/agentic-ai-threat-modeling-framework-maestro>
- **[CSA, NHI]** Cloud Security Alliance AI Safety Initiative, "The non-human identity governance vacuum: AI agents and the fastest-growing unmanaged attack surface," *CSA Research*, May 2026. [Online]. Available: <https://labs.cloudsecurityalliance.org/research/csa-whitepaper-nonhuman-identity-agentic-ai-governance-v1-cs/>
- **[Cursor incident 2025]** Cursor Community Forum, "Catastrophic damage and chaos in Plan Mode," Dec. 2025; a Cursor staff member replied in the thread that this was "a critical bug in Plan Mode constraint enforcement" and asked the user for request IDs and confirmation that Plan Mode was active. Cursor published no company postmortem. [Online]. Available: <https://forum.cursor.com/t/catastrophic-damage-and-chaos-in-plan-mode/145523>
- **[Cursor, Scaling Agents]** W. Lin (Cursor), "Scaling long-running autonomous coding," *Cursor Blog*, Jan. 2026. [Online]. Available: <https://cursor.com/blog/scaling-agents>
- **[Cursor, Self-Driving Codebases]** W. Lin (Cursor), "Towards self-driving codebases," *Cursor Blog*, Feb. 2026. [Online]. Available: <https://cursor.com/blog/self-driving-codebases>
- **[CVE-2025-32711, CVE-2025-32711]** Microsoft (CVE Numbering Authority), "CVE-2025-32711: M365 Copilot Information Disclosure Vulnerability," *CVE Program*, Jun. 2025; named EchoLeak by Aim Labs, which reported it. [Online]. Available: <https://www.cve.org/CVERecord?id=CVE-2025-32711>

- **[CVE-2025-53773, CVE-2025-53773]** Microsoft (CVE Numbering Authority), "CVE-2025-53773: GitHub Copilot and Visual Studio Remote Code Execution Vulnerability," *CVE Program*, Aug. 2025. [Online]. Available: <https://www.cve.org/CVERecord?id=CVE-2025-53773>
- **[CVE-2025-54136, CVE-2025-54136]** GitHub (CVE Numbering Authority), "CVE-2025-54136: Cursor's Modification of MCP Server Definitions Bypasses Manual Re-approvals," *CVE Program*, Aug. 2025; named MCPoison by Check Point Research, which reported it. [Online]. Available: <https://www.cve.org/CVERecord?id=CVE-2025-54136>
- **[Dang et al., Evolving Orchestration]** Y. Dang, C. Qian, X. Luo, *et al.*, "Multi-agent collaboration via evolving orchestration," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2505.19591>
- **[Das et al., Chain-of-Sanitized-Thoughts]** A. Das, S. S. Chintla, S. Endait, *et al.*, "Chain-of-Sanitized-Thoughts: Reducing PII leakage in chain-of-thought reasoning," in *Proc. ICML Workshop on AI in the Wild*, 2026. [Online]. Available: <https://arxiv.org/abs/2601.05076>
- **[Databricks, Automated Prompt Optimization]** Databricks AI Research (A. Singhvi, I. Zhou, E. Elsen, *et al.*), "Building state-of-the-art enterprise agents 90x cheaper with automated prompt optimization," *Databricks Blog*, Sep. 2025. [Online]. Available: <https://www.databricks.com/blog/building-state-art-enterprise-agents-90x-cheaper-automated-prompt-optimization>
- **[Dean & Barroso, tail at scale]** J. Dean and L. A. Barroso, "The tail at scale," *Communications of the ACM*, vol. 56, no. 2, pp. 74–80, Feb. 2013. [Online]. Available: <https://doi.org/10.1145/2408776.2408794>
- **[Debenedetti et al., AgentDojo]** E. Debenedetti, *et al.*, "AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents," in *Proc. NeurIPS*, 2024. [Online]. Available: <https://arxiv.org/abs/2406.13352>
- **[Debenedetti et al., CaMeL]** E. Debenedetti, *et al.*, "Defeating prompt injections by design (CaMeL)," *arXiv:2503.18813*, Mar. 2025. [Online]. Available: <https://arxiv.org/abs/2503.18813>
- **[Decagon, Optimizing GEPA]** R. Wang (Decagon), "Optimizing GEPA for production," *Decagon Blog*, Mar. 2026. [Online]. Available: <https://decagon.ai/blog/optimizing-gepa-for-production>

- **[Dell'Acqua et al., Navigating the jagged]** F. Dell'Acqua, E. McFowland III, E. Mollick, *et al.*, "Navigating the jagged technological frontier: Field experimental evidence of the effects of AI on knowledge worker productivity and quality," *HBS Working Paper 24-013 / SSRN 4573321*, Sep. 2023. [Online]. Available: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4573321
- **[Dell'Acqua, Falling asleep]** F. Dell'Acqua, "Falling asleep at the wheel: Human/AI collaboration in a field experiment on HR recruiters," *Laboratory for Innovation Science at Harvard, working paper*, 2022. [Online]. Available: <https://www.fabriziodellacqua.com/>
- **[Denario, Denario]** AstroPilot-AI, "Denario: Modular multi-agent system for scientific research assistance," *GitHub*, 2026. [Online]. Available: <https://github.com/AstroPilot-AI/Denario>
- **[Deng et al., ImageNet]** J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei, "ImageNet: A large-scale hierarchical image database," in *Proc. IEEE CVPR*, 2009, pp. 248–255. [Online]. Available: <https://doi.org/10.1109/CVPR.2009.5206848>
- **[Deshpande et al., TRAIL]** D. Deshpande, V. Gangal, *et al.*, "TRAIL: Trace reasoning and agentic issue localization," *arXiv:2505.08638*, May 2025. [Online]. Available: <https://arxiv.org/abs/2505.08638>
- **[Dipongkor et al., Test Coverage]** A. K. Dipongkor, T. Baral, W. Lam, and K. Moran, "Test coverage analysis of agentic pull requests," in *Proc. 42nd IEEE Int. Conf. Software Maintenance and Evolution (ICSME)*, 2026. [Online]. Available: <https://arxiv.org/abs/2607.18057>
- **[Dong et al., MINJA]** S. Dong, S. Xu, *et al.*, "MINJA: Memory injection attacks on LLM agents via query-only interaction," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2503.03704>
- **[Drouin et al., WorkArena]** A. Drouin, M. Gasse, M. Caccia, *et al.*, "WorkArena: How capable are web agents at solving common knowledge work tasks?" in *Proc. ICML*, 2024, pp. 11642–11662. [Online]. Available: <https://proceedings.mlr.press/v235/drouin24a.html>
- **[Du et al., DeepResearch Bench]** M. Du, B. Xu, C. Zhu, *et al.*, "DeepResearch Bench: A comprehensive benchmark for deep research agents," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2506.11763>
- **[Du et al., Improving factuality]** Y. Du, S. Li, A. Torralba, J. B. Tenenbaum, and I. Mordatch, "Improving factuality and reasoning in language models through

- multiagent debate," in *Proc. ICML*, 2023. [Online]. Available: <https://arxiv.org/abs/2305.14325>
- **[Ehsani et al., Agents Fail]** R. Ehsani, S. Pathak, S. Rawal, *et al.*, "Where do AI coding agents fail? An empirical study of failed agentic pull requests in GitHub," in *Proc. 23rd Int. Conf. Mining Software Repositories (MSR), Mining Challenge Track*, Apr. 2026. [Online]. Available: <https://arxiv.org/abs/2601.15195>
 - **[Ehtesham et al., survey of agent]** A. Ehtesham, *et al.*, "A survey of agent interoperability protocols: MCP, ACP, A2A, and ANP," *arXiv:2505.02279*, May 2025. [Online]. Available: <https://arxiv.org/abs/2505.02279>
 - **[El Yagoubi et al., AgentLeak]** F. El Yagoubi, G. Badu-Marfo, and R. Al Mallah, "AgentLeak: A benchmark for internal-channel privacy leakage in multi-agent LLM systems," *IEEE Access*, 2026. [Online]. Available: <https://arxiv.org/abs/2602.11510>
 - **[Erdogan et al., Plan-and-Act]** L. E. Erdogan, N. Lee, S. Kim, *et al.*, "Plan-and-Act: Improving planning of agents for long-horizon tasks," in *Proc. ICML*, 2025, pp. 15419–15462. [Online]. Available: <https://proceedings.mlr.press/v267/erdogan25a.html>
 - **[EU, AI Act]** European Parliament and Council, "Regulation (EU) 2024/1689 on artificial intelligence (EU AI Act); Articles 12, 14, 15, 50 and Annex III; high-risk-system application dates as amended by Regulation (EU) 2026/1744," *Official Journal of the European Union*, Jul. 2024. [Online]. Available: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>
 - **[EU, Digital Omnibus]** European Parliament and Council, "Regulation (EU) 2026/1744 amending Regulation (EU) 2024/1689 (Digital Omnibus on AI); standalone Annex III high-risk obligations deferred to 2 December 2027, product-embedded Annex I systems to 2 August 2028," *Official Journal of the European Union*, Jul. 2026. [Online]. Available: <https://eur-lex.europa.eu/eli/reg/2026/1744/oj>
 - **[EU, UCITS Directive]** European Parliament and Council, "Directive 2009/65/EC on the coordination of laws, regulations and administrative provisions relating to undertakings for collective investment in transferable securities (UCITS); Article 57(2)," *Official Journal of the European Union*, Jul. 2009. [Online]. Available: <https://eur-lex.europa.eu/eli/dir/2009/65/oj>
 - **[Eyuboglu et al., Cartridges]** S. Eyuboglu, R. Ehrlich, S. Arora, *et al.*, "Cartridges: Lightweight and general-purpose long context representations via self-study," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2506.06266>
 - **[FAIR Universe, Weak Lensing Challenge]** FAIR Universe Team, "Weak Lensing Uncertainty Challenge: final leaderboards," *NeurIPS 2025 Competition Track*, Lawrence

- Berkeley National Laboratory, 2025. [Online]. Available: <https://fair-universe.lbl.gov/WeakLensing-Uncertainty-Challenge.html>
- **[Fang et al., Trajectory-informed memory generation]** G. Fang, V. Isahagian, *et al.* (IBM Research), "Trajectory-informed memory generation for self-improving agent systems," *IBM Research technical report*, Mar. 2026. [Online]. Available: <https://arxiv.org/abs/2603.10600>
 - **[Federal Reserve, SR 26-2]** Board of Governors of the Federal Reserve System, Federal Deposit Insurance Corporation, and Office of the Comptroller of the Currency, "Revised guidance on model risk management (SR letter 26-2), superseding SR letter 11-7," *Federal Reserve Supervision and Regulation Letters*, Apr. 2026. [Online]. Available: <https://www.federalreserve.gov/supervisionreg/srletters/SR2602.pdf>
 - **[Feng & Pan, SentinelNet]** Y. Feng and X. Pan, "SentinelNet: Safeguarding multi-agent collaboration through credit-based dynamic threat detection," in *Proc. ACM Web Conference (WWW)*, 2026. [Online]. Available: <https://arxiv.org/abs/2510.16219>
 - **[Feng et al., GiGPO]** L. Feng, Z. Xue, T. Liu, and B. An, "Group-in-Group Policy Optimization for LLM agent training," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://openreview.net/forum?id=QXEhBMNrCW>
 - **[Feng et al., ReTool]** J. Feng, S. Huang, X. Qu, *et al.*, "ReTool: Reinforcement learning for strategic tool use in LLMs," in *Proc. ICLR*, 2026. [Online]. Available: <https://openreview.net/forum?id=tRk1nofSmz>
 - **[Ferrag et al., Protocol Exploits]** M. A. Ferrag, N. Tihanyi, D. Hamouda, L. Maglaras, A. Lakas, and M. Debbah, "From prompt injections to protocol exploits: Threats in LLM-powered AI agents workflows," *ICT Express*, vol. 12, pp. 353–383, Apr. 2026. [Online]. Available: <https://arxiv.org/abs/2506.23260>
 - **[FinOps Foundation, Token Economics]** FinOps Foundation (J. R. Storment), "Token economics: the atomic unit of AI value," *FinOps Foundation Insights*, May 2026. [Online]. Available: <https://www.finops.org/insights/token-economics-the-atomic-unit-of-ai-value/>
 - **[FINRA, Notice 24-09]** Financial Industry Regulatory Authority, "FINRA reminds members of regulatory obligations when using generative artificial intelligence and large language models (Regulatory Notice 24-09)," *FINRA Regulatory Notices*, Jun. 2024. [Online]. Available: <https://www.finra.org/rules-guidance/notices/24-09>
 - **[FINRA, Oversight Report]** Financial Industry Regulatory Authority, "2026 FINRA annual regulatory oversight report: GenAI, continuing and emerging trends," *FINRA*

Report, Dec. 2025. [Online]. Available: <https://www.finra.org/sites/default/files/2025-12/2026-annual-regulatory-oversight-report.pdf>

- **[Fish et al., Algorithmic Collusion]** S. Fish, Y. A. Gonczarowski, and R. I. Shorrer, "Algorithmic collusion by large language models," in *Proc. ACM EC*, 2026. [Online]. Available: <https://arxiv.org/abs/2404.00806>
- **[Foerster et al., COMA]** J. Foerster, G. Farquhar, T. Afouras, N. Nardelli, and S. Whiteson, "Counterfactual multi-agent policy gradients (COMA)," in *Proc. AAAI*, 2018. [Online]. Available: <https://arxiv.org/abs/1705.08926>
- **[Fourney et al., Magentic-One]** A. Fourney, *et al.*, "Magentic-One: A generalist multi-agent system for solving complex tasks," *arXiv:2411.04468*, Nov. 2024. [Online]. Available: <https://arxiv.org/abs/2411.04468>
- **[Framework observability]** AG2 / AutoGen, OpenAI Agents SDK, LangGraph + LangSmith, CrewAI, Microsoft Semantic Kernel / Microsoft Agent Framework, LlamaIndex: frameworks with built-in or pluggable OpenTelemetry observability.
- **[Froger et al., Gaia2]** R. Froger, P. Andrews, M. Bettini, *et al.*, "Gaia2: Benchmarking LLM agents on dynamic and asynchronous environments," in *Proc. ICLR*, 2026. [Online]. Available: <https://openreview.net/forum?id=9gw03JpKK4>
- **[FSB, Sound Practices]** Financial Stability Board, "Sound practices for responsible adoption of artificial intelligence (AI)," consultation report, Jun. 2026. [Online]. Available: <https://www.fsb.org/uploads/P100626.pdf>
- **[Galileo, Luna]** M. Belyi, R. Friel, S. Shao, and A. Sanyal (Galileo), "Luna: An evaluation foundation model to catch language model hallucinations with high accuracy and low cost," *arXiv:2406.00975*, Jun. 2024. [Online]. Available: <https://arxiv.org/abs/2406.00975>
- **[Galileo, Luna-2]** V. Goel, R. Dsouza, N. Ega, *et al.* (Galileo), "Luna-2: Scalable single-token evaluation with small language models," *arXiv:2602.18583*, Feb. 2026. [Online]. Available: <https://arxiv.org/abs/2602.18583>
- **[Galileo, OWASP ASI02]** P. Bhavsar, M. Branconier, O. Adoh, and P. Lacey (Galileo), "OWASP ASI02: When agents weaponize their own tools," *Galileo Blog*, May 2026. [Online]. Available: <https://galileo.ai/blog/owasp-agentic-ai-asi02-tool-misuse>
- **[Galster et al., Configuring Agentic Tools]** M. Galster, S. Mohsenimofidi, J. L. Lulla, *et al.*, "Configuring agentic AI coding tools: An exploratory study," in *Proc. 3rd ACM Int.*

- Conf. AI-powered Software (AIware)*, Jul. 2026. [Online]. Available: <https://arxiv.org/abs/2602.14690>
- **[Gamma et al., Design Patterns]** E. Gamma, R. Helm, R. Johnson, and J. Vlissides, *Design Patterns: Elements of Reusable Object-Oriented Software*. Reading, MA: Addison-Wesley, 1994. [Online]. Available: <https://www.pearson.com/en-us/subject-catalog/p/design-patterns-elements-of-reusable-object-oriented-software/P200000009480>
 - **[Gao et al., survey of self-evolving]** H.-A. Gao, *et al.*, "A survey of self-evolving agents: What, when, how, and where to evolve on the path to artificial super intelligence," *Transactions on Machine Learning Research (TMLR)*, 2026. [Online]. Available: <https://arxiv.org/abs/2507.21046>
 - **[Gartner, Enterprise Adoption]** Gartner, "Gartner predicts 40% of enterprise apps will feature task-specific AI agents by 2026, up from less than 5% in 2025," *Gartner press release*, Aug. 2025. [Online]. Available: <https://www.gartner.com/en/newsroom/press-releases/2025-08-26-gartner-predicts-40-percent-of-enterprise-apps-will-feature-task-specific-ai-agents-by-2026-up-from-less-than-5-percent-in-2025>
 - **[GCP Secret Manager, Secret Manager]** Google Cloud, "Secret Manager — managed storage and access for API keys, passwords, and certificates," *Google Cloud Documentation*, continuously updated. [Online]. Available: <https://docs.cloud.google.com/secret-manager/docs>
 - **[GCP Tasks, Cloud Tasks]** Google Cloud, "Cloud Tasks — managed asynchronous task execution with at-least-once delivery," *Google Cloud Documentation*, continuously updated. [Online]. Available: <https://docs.cloud.google.com/tasks/docs>
 - **[GDPR, EU]** European Parliament and Council, "Regulation (EU) 2016/679: General Data Protection Regulation (Article 17 — right to erasure)," *Official Journal of the European Union*, Apr. 2016. [Online]. Available: <https://eur-lex.europa.eu/eli/reg/2016/679/oj>
 - **[Gehring et al., RLEF]** J. Gehring, K. Zheng, J. Copet, *et al.*, "RLEF: Grounding code LLMs in execution feedback with reinforcement learning," in *Proc. ICML*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.02089>
 - **[Geifman & El-Yaniv, Selective classification]** Y. Geifman and R. El-Yaniv, "Selective classification for deep neural networks," in *Proc. NeurIPS*, 2017. [Online]. Available: <https://arxiv.org/abs/1705.08500>

- **[Geng & Neubig, Asynchronous SWE agents]** J. Geng and G. Neubig, "Effective strategies for asynchronous software engineering agents," in *Proc. COLM*, 2026. [Online]. Available: <https://arxiv.org/abs/2603.21489>
- **[GitHub, Copilot Usage-Based Billing]** GitHub, "GitHub Copilot is moving to usage-based billing," *The GitHub Blog*, Apr. 2026. [Online]. Available: <https://github.blog/news-insights/company-news/github-copilot-is-moving-to-usage-based-billing/>
- **[Goodhart, Goodhart's Law]** C. A. E. Goodhart, "Problems of monetary management: The U.K. experience," *Papers in Monetary Economics*, vol. 1, Reserve Bank of Australia, 1975; reprinted in C. A. E. Goodhart, *Monetary Theory and Practice: The UK Experience*, London: Macmillan, 1984, pp. 91–121. [Online]. Available: https://doi.org/10.1007/978-1-349-17295-5_4
- **[Google Cloud Run]** Google Cloud, "Cloud Run: serverless containers for stateless workloads," *Google Cloud Documentation*, continuously updated. [Online]. Available: <https://docs.cloud.google.com/run/docs>
- **[Google SRE Workbook]** B. Beyer, et al., Eds., *The Site Reliability Workbook: Practical Ways to Implement SRE* (Chapter 5: Alerting on SLOs). Sebastopol, CA: O'Reilly Media, 2018. [Online]. Available: <https://sre.google/workbook/alerting-on-slos/>
- **[Google, A2A]** Google → Linux Foundation, "Agent-to-Agent (A2A) Protocol specification," *a2a-protocol.org*, Apr. 2025. [Online]. Available: <https://a2a-protocol.org/latest/>
- **[Google, A2UI]** Google A2UI Team, "Introducing A2UI: An open project for agent-driven interfaces," *Google Developers Blog*, Dec. 2025. [Online]. Available: <https://developers.googleblog.com/introducing-a2ui-an-open-project-for-agent-driven-interfaces/>
- **[Google, ADK]** Google, "Agent Development Kit (ADK) documentation," *Google Cloud Docs*, 2025. [Online]. Available: <https://adk.dev/>
- **[Google, Agent Sandbox]** Google Open Source / Kubernetes SIG Apps, "Unleashing autonomous AI agents: Why Kubernetes needs a new standard for agent execution (Agent Sandbox)," *Google Open Source Blog*, Nov. 2025. [Online]. Available: <https://opensource.googleblog.com/2025/11/unleashing-autonomous-ai-agents-why-kubernetes-needs-a-new-standard-for-agent-execution.html>
- **[Google, AP2]** Google, "Agent Payments Protocol (AP2) v0.2 — Human Not Present payments; donated to the FIDO Alliance," *Google Blog*, Apr. 2026. [Online]. Available:

<https://blog.google/products-and-platforms/platforms/google-pay/agent-payments-protocol-fido-alliance/>

- **[Google, Gemini Pricing]** Google, "Gemini API pricing — per-million-token rates by model, including context-caching rates," *Google AI for Developers*, continuously updated. [Online]. Available: <https://ai.google.dev/gemini-api/docs/pricing>
- **[Gou et al., CRITIC]** Z. Gou, Z. Shao, Y. Gong, *et al.*, "CRITIC: Large language models can self-correct with tool-interactive critiquing," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2305.11738>
- **[Gou et al., Mind2Web 2]** B. Gou, Z. Huang, Y. Ning, *et al.*, "Mind2Web 2: Evaluating agentic search with Agent-as-a-Judge," in *Proc. NeurIPS Datasets and Benchmarks Track*, 2025. [Online]. Available: <https://arxiv.org/abs/2506.21506>
- **[Green et al., Leaky thoughts]** T. Green, M. Gubri, *et al.*, "Leaky thoughts: Large reasoning models are not private thinkers," in *Proc. EMNLP*, 2025. [Online]. Available: <https://arxiv.org/abs/2506.15674>
- **[Greenblatt et al., Independent investigation]** R. Greenblatt, A. Cotra, and H. Wijk (METR and Redwood Research), "Brief independent investigation of agents' behavior, reasoning and collaboration in the OpenAI / Hugging Face hacking incident," *METR report*, Aug. 2026. [Online]. Available: <https://metr.org/blog/2026-08-26-openai-hugging-face-incident-investigation/>
- **[Greshake et al., Not what you've]** K. Greshake, *et al.*, "Not what you've signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection," in *Proc. AISEC @ CCS*, 2023. [Online]. Available: <https://arxiv.org/abs/2302.12173>
- **[Grötschla et al., AgentsNet]** F. Grötschla, *et al.*, "AgentsNet: Coordination and collaborative reasoning in multi-agent LLMs," *arXiv:2507.08616*, Jul. 2025. [Online]. Available: <https://arxiv.org/abs/2507.08616>
- **[Gu et al., Agent Smith]** X. Gu, *et al.*, "Agent Smith: A single image can jailbreak one million multimodal LLM agents exponentially fast," in *Proc. ICML*, 2024. [Online]. Available: <https://arxiv.org/abs/2402.08567>
- **[Guo et al., Large language model]** T. Guo, X. Chen, Y. Wang, *et al.*, "Large language model based multi-agents: A survey of progress and challenges," in *Proc. IJCAI*, 2024. [Online]. Available: <https://arxiv.org/abs/2402.01680>
- **[Gupta et al., Support at Scale]** A. Gupta, K. R. Rossell, E. Alcobaça, *et al.* (Nubank), "Building customer support AI agents at 100M-user scale: An evaluation-driven

- framework," in *Proc. 32nd ACM SIGKDD Conf. Knowledge Discovery and Data Mining (KDD '26)*, Aug. 2026. [Online]. Available: <https://arxiv.org/abs/2606.08867>
- **[Harbach and Johnson, Surface Attribution]** M. Harbach and J. Johnson, "Who am I talking to? A large-scale measurement of surface attribution across real-world security and privacy interfaces," in *Proc. ACM CHI*, 2026. [Online]. Available: <https://research.google/pubs/who-am-i-talking-to-a-large-scale-measurement-of-surface-attribution-across-real-world-security-and-privacy-interfaces/>
 - **[HashiCorp Vault, Vault]** HashiCorp, "Vault — identity-based secrets and encryption management," *HashiCorp Documentation*, continuously updated. [Online]. Available: <https://developer.hashicorp.com/vault/docs>
 - **[He et al., Agent-in-the-Middle]** P. He, *et al.*, "Red-teaming LLM multi-agent systems via communication attacks (Agent-in-the-Middle)," in *Proc. ACL*, 2025. [Online]. Available: <https://arxiv.org/abs/2502.14847>
 - **[He et al., Cursor Velocity]** H. He, C. Miller, S. Agarwal, *et al.*, "Speed at the cost of quality: How Cursor AI increases short-term velocity and long-term complexity in open-source projects," in *Proc. 23rd Int. Conf. Mining Software Repositories (MSR)*, Apr. 2026, pp. 181–193. [Online]. Available: <https://arxiv.org/abs/2511.04427>
 - **[Hermes, Hermes Agent]** Nous Research, "Hermes Agent: a self-improving personal AI agent with persistent memory," *GitHub*, 2026. [Online]. Available: <https://github.com/NousResearch/hermes-agent>
 - **[HHS, HIPAA Medical Records Retention]** U.S. Department of Health and Human Services, "Does the HIPAA Privacy Rule require covered entities to keep patients' medical records for any period of time?" Jul. 26, 2013. [Online]. Available: <https://www.hhs.gov/hipaa/for-professionals/faq/580/does-hipaa-require-covered-entities-to-keep-medical-records-for-any-period/index.html>
 - **[HHS, HIPAA Security Documentation]** U.S. Department of Health and Human Services, "Summary of the HIPAA Security Rule — Policies and Procedures and Documentation Requirements," continuously updated. [Online]. Available: <https://www.hhs.gov/hipaa/for-professionals/security/laws-regulations/index.html>
 - **[Ho & Ren, Cell Layout]** C.-T. Ho and H. Ren, "Large language model (LLM) for standard cell layout design optimization," in *Proc. IEEE LLM Aided Design Workshop (LAD)*, 2024. [Online]. Available: <https://arxiv.org/abs/2406.06549>
 - **[Ho et al., Marco]** C.-T. Ho, J. Gong, Y. Bai, C. Deng, H. Ren, and B. Khailany, "Marco: Configurable graph-based task solving and multi-AI agents framework for hardware

- design," in *Proc. IEEE Symp. VLSI Technology and Circuits*, Jun. 2025. [Online]. Available: <https://arxiv.org/abs/2504.01962>
- **[Ho et al., Polymath]** C.-T. Ho, J. Gong, H. Yang, A. B. Akkur, and H. Ren, "Polymath: Self-improving hierarchical workflow for multi-domain problem solving," in *Proc. ACM Int. Symp. Physical Design (ISPD)*, Mar. 2026. [Online]. Available: <https://doi.org/10.1145/3764386.3779639>
 - **[Ho et al., VerilogCoder]** C.-T. Ho, H. Ren, and B. Khailany, "VerilogCoder: Autonomous Verilog coding agents with graph-based planning and abstract syntax tree (AST)-based waveform tracing tool," in *Proc. AAAI Conf. Artificial Intelligence*, Feb. 2025, pp. 300–307. [Online]. Available: <https://arxiv.org/abs/2408.08927>
 - **[Hong et al., Context Rot]** K. Hong, A. Troynikov, and J. Huber, "Context Rot: How increasing input tokens impacts LLM performance," *Chroma Research*, Jul. 2025. [Online]. Available: <https://www.trychroma.com/research/context-rot>
 - **[Hong et al., MetaGPT]** S. Hong, M. Zhuge, J. Chen, X. Zheng, *et al.*, "MetaGPT: Meta programming for a multi-agent collaborative framework," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2308.00352>
 - **[Hu et al., MemoryAgentBench]** Y. Hu, Y. Wang, and J. McAuley, "Evaluating memory in LLM agents via incremental multi-turn interactions (MemoryAgentBench)," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2507.05257>
 - **[Huang et al., Cannot Self-Correct]** J. Huang, *et al.*, "Large language models cannot self-correct reasoning yet," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2310.01798>
 - **[Huang et al., MetaTool]** Y. Huang, *et al.*, "MetaTool benchmark for large language models: Deciding whether to use tools and which to use," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2310.03128>
 - **[Huang et al., MAgentBench]** Q. Huang, J. Vora, P. Liang, and J. Leskovec, "MAgentBench: Evaluating language agents on machine learning experimentation," in *Proc. ICML*, 2024. [Online]. Available: <https://arxiv.org/abs/2310.03302>
 - **[Hubinger et al., Sleeper agents]** E. Hubinger, *et al.*, "Sleeper agents: Training deceptive LLMs that persist through safety training," *arXiv:2401.05566*, Jan. 2024. [Online]. Available: <https://arxiv.org/abs/2401.05566>
 - **[HubSpot, Outcome Pricing]** HubSpot, "HubSpot's Customer Agent and Prospecting Agent: Now you pay when the task is complete," *HubSpot Company News*, Apr. 2026.

- [Online]. Available: <https://www.hubspot.com/company-news/hubspots-customer-agent-and-prospecting-agent-now-you-pay-when-the-task-is-complete>
- **[Hugging Face, OpenEnv]** Hugging Face, "OpenEnv: Agentic execution environments," *GitHub*, continuously updated; checked Sep. 2026. [Online]. Available: <https://github.com/huggingface/OpenEnv>
 - **[Hugging Face, Pickle]** Hugging Face, "Pickle scanning and model security on the Hub," *Hugging Face Documentation*, 2025. [Online]. Available: <https://huggingface.co/docs/hub/en/security-pickle>
 - **[Hugging Face, Security incident]** Hugging Face, "Security incident disclosure — July 2026," *Hugging Face Blog*, Jul. 2026. [Online]. Available: <https://huggingface.co/blog/security-incident-july-2026>
 - **[Hugging Face, smolagents]** Hugging Face, "smolagents: A minimal library to build code-driven AI agents," *GitHub*, 2024. [Online]. Available: <https://github.com/huggingface/smolagents>
 - **[Huo et al., BIRD-INTERACT]** N. Huo, X. Xu, J. Li, *et al.*, "BIRD-INTERACT: Re-imagining text-to-SQL evaluation via lens of dynamic interactions," in *Proc. ICLR*, 2026. [Online]. Available: <https://openreview.net/forum?id=nHrYBGujps>
 - **[IBM, ACP]** IBM / BeeAI, "Agent Communication Protocol (ACP) specification (repository archived Aug. 2025; protocol merged into A2A; see [LF AI & Data, ACP Joins A2A])," 2025. [Online]. Available:
 - **[IBM, Granite Guardian 3.3]** IBM, "Granite Guardian 3.3: Open-weight safety model," *IBM Product Documentation*, 2025. [Online]. Available: <https://www.ibm.com/granite/docs/models/guardian>
 - **[IBM, ITBench MAST]** IBM Research and UC Berkeley, "IBM and UC Berkeley diagnose why enterprise agents fail using IT-Bench and MAST," *IBM Research Blog (Hugging Face)*, Feb. 2026. [Online]. Available: <https://huggingface.co/blog/ibm-research/itbenchandmast>
 - **[Intercom, Fin Outcomes]** Intercom, "Fin AI Agent outcomes," *Intercom Help Center*, continuously updated; checked Sep. 2026. [Online]. Available: <https://www.intercom.com/help/en/articles/8205718-fin-ai-agent-outcomes>
 - **[Invariant Labs, MCP security notification]** Invariant Labs, "MCP security notification: tool poisoning attacks, rug pulls, and tool shadowing," *Invariant Labs Blog*, Apr. 2025. [Online]. Available: <https://invariantlabs.ai/blog/mcp-security-notification-tool-poisoning-attacks>

- **[Istio Docs, Istio service mesh]** Istio Authors, "Istio service mesh documentation (mTLS, PeerAuthentication, traffic management)," *istio.io*, continuously updated. [Online]. Available: <https://istio.io/latest/docs/>
- **[JFrog, Data scientists targeted]** JFrog Security Research, "Data scientists targeted by malicious Hugging Face ML models with silent backdoor," *JFrog Blog*, Feb. 2024. [Online]. Available: <https://jfrog.com/blog/data-scientists-targeted-by-malicious-hugging-face-ml-models-with-silent-backdoor/>
- **[Jha et al., ControlValve]** R. Jha, H. Triedman, J. Wagle, and V. Shmatikov, "Breaking and fixing defenses against control-flow hijacking in multi-agent systems (ControlValve)," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2510.17276>
- **[Jha et al., ITBench]** S. Jha, R. R. Arora, Y. Watanabe, *et al.*, "ITBench: Evaluating AI agents across diverse real-world IT automation tasks," in *Proc. ICML*, 2025, pp. 27134–27197. [Online]. Available: <https://proceedings.mlr.press/v267/jha25a.html>
- **[Jimenez et al., SWE-bench]** C. E. Jimenez, J. Yang, A. Wettig, *et al.*, "SWE-bench: Can language models resolve real-world GitHub issues?" in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2310.06770>
- **[JSON-RPC, Specification]** JSON-RPC Working Group, "JSON-RPC 2.0 specification," *jsonrpc.org*, 2010. [Online]. Available: <https://www.jsonrpc.org/specification>
- **[Jwalapuram et al., Illusion]** P. Jwalapuram, *et al.*, "The illusion of multi-agent advantage," *arXiv:2606.13003*, Jun. 2026. [Online]. Available: <https://arxiv.org/abs/2606.13003>
- **[Kafka Docs, Apache Kafka]** Apache Software Foundation, "Apache Kafka — distributed event streaming platform," *kafka.apache.org*, continuously updated. [Online]. Available: <https://kafka.apache.org/documentation/>
- **[Kahlon et al., Agent-Initiated Interaction]** N. Kahlon, G. Rom, T. Efron, *et al.*, "Agent-initiated interaction in phone UI automation," in *Companion Proc. ACM Web Conf. (WWW '25)*, May 2025, pp. 2391–2400. [Online]. Available: <https://doi.org/10.1145/3701716.3717526>
- **[Kapoor et al., AI agents that]** S. Kapoor, B. Stroebel, Z. S. Siegel, N. Nadgir, and A. Narayanan, "AI agents that matter," *Trans. Mach. Learn. Res.*, 2024. [Online]. Available: <https://arxiv.org/abs/2407.01502>

- **[Kapoor et al., Holistic Leaderboard]** S. Kapoor, *et al.*, "Holistic agent leaderboard: The missing infrastructure for AI agent evaluation," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2510.11977>
- **[Kartik et al., AgentCompass]** N. V. J. K. Kartik, G. Sapra, *et al.*, "AgentCompass: Towards reliable evaluation of agentic workflows in production," *arXiv:2509.14647*, Sep. 2025. [Online]. Available: <https://arxiv.org/abs/2509.14647>
- **[Kavathekar et al., TAMAS]** I. Kavathekar, H. Jain, A. Rathod, P. Kumaraguru, and T. Ganu, "TAMAS: Benchmarking adversarial risks in multi-agent LLM systems," in *Proc. ACL (Volume 1: Long Papers)*, 2026, pp. 31238–31268. [Online]. Available: <https://arxiv.org/abs/2511.05269>
- **[Ke et al., CritiqueLLM]** P. Ke, *et al.*, "CritiqueLLM: Towards an informative critique generation model for evaluation of large language model generation," in *Proc. ACL*, 2024, pp. 13034–13054. [Online]. Available: <https://arxiv.org/abs/2311.18702>
- **[Kephart & Chess, Autonomic Vision]** J. O. Kephart and D. M. Chess, "The vision of autonomic computing," *Computer*, vol. 36, no. 1, pp. 41–50, Jan. 2003. [Online]. Available: <https://doi.org/10.1109/MC.2003.1160055>
- **[Khalifa et al., Gaming the Judge]** M. Khalifa, L. Logeswaran, J. Kim, *et al.*, "Gaming the judge: Unfaithful chain-of-thought can undermine agent evaluation," *arXiv:2601.14691*, Jan. 2026. [Online]. Available: <https://arxiv.org/abs/2601.14691>
- **[Khattab et al., DSPy]** O. Khattab, *et al.*, "DSPy: Compiling declarative language model calls into self-improving pipelines," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2310.03714>
- **[Kim et al., Correlated Errors]** E. Kim, A. Garg, K. Peng, and N. Garg, "Correlated errors in large language models," in *Proc. ICML*, 2025. [Online]. Available: <https://arxiv.org/abs/2506.07962>
- **[Kim et al., Prometheus 2]** S. Kim, *et al.*, "Prometheus 2: An open source language model specialized in evaluating other language models," in *Proc. EMNLP*, 2024. [Online]. Available: <https://arxiv.org/abs/2405.01535>
- **[Kim et al., Scaling Agent Systems]** Y. Kim, *et al.*, "Towards a science of scaling agent systems," *arXiv:2512.08296*, v3, Apr. 2026. [Online]. Available: <https://arxiv.org/abs/2512.08296>
- **[Kimi Team, Kimi K2.5]** Kimi Team (Moonshot AI), "Kimi K2.5: Visual agentic intelligence," *arXiv:2602.02276*, technical report, Feb. 2026. [Online]. Available: <https://arxiv.org/abs/2602.02276>

- **[Kiro incident 2025]** B. Pariseau, "AWS Kiro 'user error' reflects common AI coding review gap," *TechTarget*, 23 Feb. 2026, reporting the Dec. 2025 Cost Explorer outage and Amazon's 20 Feb. 2026 statement, which attributed it to misconfigured access controls rather than to the agent. The thirteen-hour duration comes from Financial Times reporting via AI Incident Database 1442, not from this article, and so do the permission details the book marks "as reported": the Financial Times (Feb. 2026, relayed by *The Decoder*) reported that Kiro was given the operator's own permissions and that no second person's approval was needed, "something that would normally be required." [Online]. Available: <https://www.techtarget.com/it-infrastructure/news/366639129/AWS-Kiro-user-error-reflects-common-AI-coding-review-gap>
- **[Koi Security, ClawHavoc]** Koi Security, "ClawHavoc: 341 malicious ClawedBot skills found by the bot they were targeting," *Koi Security Blog*, Feb. 2026, cited from the Internet Archive's 6 March 2026 snapshot because koi.ai was folded into Palo Alto Networks after an acquisition and every path on the domain, this post included, now redirects to a Cortex product page. [Online]. Available: <https://web.archive.org/web/20260306131006/https://www.koi.ai/blog/clawhavoc-341-malicious-clawedbot-skills-found-by-the-bot-they-were-targeting>
- **[Krakovna et al., Specification Gaming]** V. Krakovna, J. Uesato, V. Mikulik, *et al.*, "Specification gaming: The flip side of AI ingenuity," *DeepMind Blog*, Apr. 2020. [Online]. Available: <https://www.deepmind.com/blog/specification-gaming-the-flip-side-of-ai-ingenuity>
- **[Krizhevsky et al., AlexNet]** A. Krizhevsky, I. Sutskever, and G. E. Hinton, "ImageNet classification with deep convolutional neural networks (AlexNet)," in *Proc. NeurIPS*, 2012. [Online]. Available: <https://papers.nips.cc/paper/2012/hash/c399862d3b9d6b76c8436e924a68c45b-Abstract.html>
- **[Kubernetes Docs, Kubernetes documentation]** CNCF / Kubernetes Authors, "Kubernetes documentation (Deployments, HPA, NetworkPolicy, Secrets, terminationGracePeriodSeconds)," *kubernetes.io*, continuously updated. [Online]. Available: <https://kubernetes.io/docs/>
- **[Kumar et al., SCoRe]** A. Kumar, V. Zhuang, R. Agarwal, Y. Su, *et al.*, "Training language models to self-correct via reinforcement learning (SCoRe)," in *Proc. ICLR Oral*, 2025. [Online]. Available: <https://openreview.net/forum?id=CjwERcAU7w>
- **[Kumar, OAuth Agent Grants]** S. Kumar, "OAuth profile for delegated AI agent authorization," *IETF Internet-Draft draft-mishra-oauth-agent-grants-02*, Aug. 2026,

- work in progress. [Online]. Available: <https://datatracker.ietf.org/doc/draft-mishra-oauth-agent-grants/>
- **[Kuntz et al., OS-Harm]** T. Kuntz, A. Duzan, H. Zhao, *et al.*, "OS-Harm: A benchmark for measuring safety of computer use agents," in *Proc. NeurIPS Datasets and Benchmarks*, 2025. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2025/hash/4009bfff0cd87ba2203c8e3a2f082aaec-Abstract-Datasets_and_Benchmarks_Track.html
 - **[Kuratov et al., BABILong]** Y. Kuratov, A. Bulatov, P. Anokhin, *et al.*, "BABILong: Testing the limits of LLMs with long context reasoning-in-a-haystack," in *Proc. NeurIPS*, 2024. [Online]. Available: <https://arxiv.org/abs/2406.10149>
 - **[Kwon et al., PagedAttention]** W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. E. Gonzalez, H. Zhang, and I. Stoica, "Efficient memory management for large language model serving with PagedAttention (vLLM)," in *Proc. SOSP*, 2023. [Online]. Available: <https://arxiv.org/abs/2309.06180>
 - **[Lakera, Lakera Guard]** Lakera (a Check Point company), "Lakera Guard: real-time AI firewall for prompt injection, data leakage, and content moderation (vendor-stated sub-50ms runtime latency)," *lakera.ai*, 2025. [Online]. Available: <https://www.lakera.ai/lakera-guard>
 - **[Lakshmanan, EchoLeak]** R. Lakshmanan, "Zero-click AI vulnerability exposes Microsoft 365 Copilot data without user interaction," *The Hacker News*, 12 Jun. 2025, reporting Aim Labs' EchoLeak research: the attacker's email is retrieved when the user asks Copilot a question, and data leaves "via Microsoft Teams and SharePoint URLs"; "no user clicks are required." [Online]. Available: <https://thehackernews.com/2025/06/zero-click-ai-vulnerability-exposes.html>
 - **[Lam et al., Just-In-Time Objectives]** M. S. Lam, O. Shaikh, H. Xu, *et al.*, "Just-in-time objectives: A general approach for specialized AI interactions," in *Proc. ACM CHI*, 2026. [Online]. Available: <https://doi.org/10.1145/3772318.3790713>
 - **[Landis & Koch, measurement of observer]** J. R. Landis and G. G. Koch, "The measurement of observer agreement for categorical data," *Biometrics*, vol. 33, no. 1, pp. 159–174, Mar. 1977. [Online]. Available: <https://doi.org/10.2307/2529310>
 - **[LangChain, Agent Engineering]** LangChain, "State of agent engineering: survey of 1,340 teams on agent adoption, observability, and evaluation practice (fielded Nov. 18–Dec. 2, 2025)," *LangChain*, Dec. 2025. [Online]. Available: <https://www.langchain.com/state-of-agent-engineering>

- **[LangChain, LangGraph]** LangChain, "LangGraph: Graph-based stateful agent workflows," *GitHub*, 2024. [Online]. Available: <https://github.com/langchain-ai/langgraph>
- **[LangChain, LangGraph Interrupts]** LangChain, "LangGraph interrupts — pausing a graph on interrupt and resuming from the persisted state," *LangChain Documentation*, continuously updated. [Online]. Available: <https://docs.langchain.com/oss/python/langgraph/interrupts>
- **[LangChain, Memory]** LangChain, "Short-term memory: conversation history, trimming, deletion, and summarization in LangChain agents," *LangChain Documentation*, continuously updated. [Online]. Available: <https://docs.langchain.com/oss/python/langchain/short-term-memory>
- **[LangChain, Persistence]** LangChain, "LangGraph persistence — checkpoints, threads, and durable state," *LangChain Documentation*, continuously updated. [Online]. Available: <https://docs.langchain.com/oss/python/langgraph/persistence>
- **[LangSmith Cloud, LangGraph Platform]** LangChain, "LangSmith Deployment (formerly LangGraph Platform) — managed deployment, tracing, and evaluation for LangGraph agents," *LangChain Documentation*, continuously updated. [Online]. Available: <https://docs.langchain.com/langsmith/deployment>
- **[Lee & Tiwari, Prompt Infection]** D. Lee, M. Tiwari, and B. Miranda, "Prompt Infection: LLM-to-LLM prompt injection within multi-agent systems," in *Proc. ESORICS 2025 International Workshops*, LNCS, Springer, 2026, pp. 511–520. [Online]. Available: <https://arxiv.org/abs/2410.07283>
- **[Lei et al., Spider 2.0]** F. Lei, J. Chen, *et al.*, "Spider 2.0: Evaluating language models on real-world enterprise text-to-SQL workflows," in *Proc. ICLR*, 2025. [Online]. Available: <https://openreview.net/forum?id=XmProj9cPs>
- **[Leviathan et al., Generative UI]** Y. Leviathan, D. Valevski, *et al.* (Google Research), "Generative UI: LLMs are effective UI generators," *Google Research*, Apr. 2026. [Online]. Available: <https://arxiv.org/abs/2604.09577>
- **[Leviathan et al., Speculative Decoding]** Y. Leviathan, M. Kalman, and Y. Matias, "Fast inference from transformers via speculative decoding," in *Proc. ICML*, 2023. [Online]. Available: <https://arxiv.org/abs/2211.17192>
- **[LF AI & Data, ACP Joins A2A]** K. Blair and T. Segal, "ACP joins forces with A2A under the Linux Foundation's LF AI & Data," *LF AI & Data Community Blog*, Aug. 2025.

- [Online]. Available: <https://lfaidata.foundation/communityblog/2025/08/29/acp-joins-forces-with-a2a-under-the-linux-foundations-lf-ai-data/>
- **[Li et al., AgentFlow]** Z. Li, H. Zhang, S. Han, *et al.*, "In-the-flow agentic system optimization for effective planning and tool use," in *Proc. ICLR*, 2026. [Online]. Available: <https://openreview.net/forum?id=Mf5A1eTUVK>
 - **[Li et al., AIDev]** H. Li, H. Zhang, and A. E. Hassan, "AIDev: Studying AI coding agents on GitHub," in *Proc. 23rd Int. Conf. Mining Software Repositories (MSR)*, Apr. 2026. [Online]. Available: <https://arxiv.org/abs/2602.09185>
 - **[Li et al., Auto-J]** J. Li, *et al.*, "Generative judge for evaluating alignment," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2310.05470>
 - **[Li et al., Embodied Agent Interface]** M. Li, S. Zhao, Q. Wang, *et al.*, "Embodied Agent Interface: Benchmarking LLMs for embodied decision making," in *Proc. NeurIPS Datasets and Benchmarks Track*, 2024. [Online]. Available: <https://arxiv.org/abs/2410.07166>
 - **[Li et al., Reflection-Bench]** L. Li, *et al.*, "Reflection-Bench: Evaluating epistemic agency in large language models," in *Proc. ICML*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.16270>
 - **[Li et al., SkillsBench]** X. Li, *et al.*, "SkillsBench: Benchmarking how well agent skills work across diverse tasks," *arXiv:2602.12670*, Feb. 2026. [Online]. Available: <https://arxiv.org/abs/2602.12670>
 - **[Li et al., WebDevJudge]** C. Li, Y. Zheng, *et al.*, "WebDevJudge: Evaluating (M)LLMs as critiques for web development quality," in *Proc. ICLR*, 2026. [Online]. Available: <https://openreview.net/forum?id=CCSPm6V5EF>
 - **[Liang et al., HELM]** P. Liang, R. Bommasani, T. Lee, *et al.*, "Holistic evaluation of language models (HELM)," *Trans. Mach. Learn. Res.*, 2023. [Online]. Available: <https://arxiv.org/abs/2211.09110>
 - **[Liang et al., MAD Debate]** T. Liang, Z. He, W. Jiao, *et al.*, "Encouraging divergent thinking in large language models through multi-agent debate," in *Proc. EMNLP*, 2023. [Online]. Available: <https://arxiv.org/abs/2305.19118>
 - **[Liao et al., RedTeamCUA]** Z. Liao, *et al.*, "RedTeamCUA: Realistic adversarial testing of computer-use agents in hybrid web-OS environments," in *Proc. ICLR*, 2026 (ICLR 2026). [Online]. Available: <https://arxiv.org/abs/2505.21936>

- **[Lightman et al., Verify Step-by-Step]** H. Lightman, V. Kosaraju, Y. Burda, *et al.*, "Let's verify step by step," in *Proc. ICLR*, 2023. [Online]. Available: <https://arxiv.org/abs/2305.20050>
- **[Lin et al., BAGEN]** Y. Lin, *et al.*, "Towards budget-aware agents: Do LLM agents know what they will spend?," in *Proc. ICML Workshop on AI in the Wild*, 2026. [Online]. Available: <https://arxiv.org/abs/2606.00198>
- **[Lin et al., Sleep-time compute]** K. Lin, *et al.*, "Sleep-time compute: Beyond inference scaling at test-time," *arXiv:2504.13171*, Apr. 2025. [Online]. Available: <https://arxiv.org/abs/2504.13171>
- **[Linkerd Docs, mTLS, policy, observability]** Linkerd Authors, "Linkerd service mesh documentation (mTLS, policy, observability)," *linkerd.io*, continuously updated. [Online]. Available: <https://linkerd.io/docs/>
- **[Linux Foundation, A2A First Year]** Linux Foundation, "A2A Protocol surpasses 150 organizations, lands in major cloud platforms, and sees enterprise production use in first year," *Linux Foundation Press*, Apr. 2026. [Online]. Available: <https://www.linuxfoundation.org/press/a2a-protocol-surpasses-150-organizations-lands-in-major-cloud-platforms-and-sees-enterprise-production-use-in-first-year>
- **[Linux Foundation, AAIIF]** Linux Foundation, "Linux Foundation announces the formation of the Agentic AI Foundation (AAIF), anchored by new project contributions including Model Context Protocol (MCP), goose, and AGENTS.md," *Linux Foundation Press*, Dec. 2025. [Online]. Available: <https://www.linuxfoundation.org/press/linux-foundation-announces-the-formation-of-the-agentic-ai-foundation>
- **[Linux Foundation, Tokenomics Foundation]** Linux Foundation, "Linux Foundation launches the Tokenomics Foundation to define the economics and ROI of AI value," *Linux Foundation News*, Aug. 4, 2026. [Online]. Available: <https://www.linuxfoundation.org/press/linux-foundation-launches-the-tokenomics-foundation-to-define-the-economics-and-roi-of-ai-value>
- **[Linux Foundation, x402 Foundation]** Linux Foundation, "Linux Foundation announces operational launch of x402 Foundation to standardize internet-native payments for AI agents and applications," *Linux Foundation Press*, Jul. 2026. [Online]. Available: <https://www.linuxfoundation.org/press/linux-foundation-announces-operational-launch-of-x402-foundation-to-standardize-internet-native-payments-for-ai-agents-and-applications>

- **[LiteLLM, LiteLLM]** BerriAI, "LiteLLM — open-source LLM gateway with a unified API, routing, rate limiting, and budgets across 100+ providers," *LiteLLM Documentation*, continuously updated. [Online]. Available: <https://docs.litellm.ai/>
- **[LiteLLM, Security Update]** BerriAI, "Security Update: Suspected Supply Chain Incident," *LiteLLM Documentation*, Mar. 2026. [Online]. Available: <https://docs.litellm.ai/blog/security-update-march-2026>
- **[Liu et al., AgentBench]** X. Liu, *et al.*, "AgentBench: Evaluating LLMs as agents," in *Proc. ICLR*, 2024. [Online]. Available: <https://openreview.net/forum?id=zAdUB0aCTQ>
- **[Liu et al., AutoDAN-Turbo]** X. Liu, *et al.*, "AutoDAN-Turbo: A lifelong agent for strategy self-exploration to jailbreak LLMs," in *Proc. ICLR Spotlight*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.05295>
- **[Liu et al., Dive into Claude Code]** J. Liu, X. Zhao, X. Shang, and Z. Shen, "Dive into Claude Code: The design space of today's and future AI agent systems," *arXiv: 2604.14228*, Jul. 2026. [Online]. Available: <https://arxiv.org/abs/2604.14228>
- **[Liu et al., Lost in the Middle]** N. F. Liu, K. Lin, J. Hewitt, *et al.*, "Lost in the middle: How language models use long contexts," *Transactions of the Association for Computational Linguistics*, vol. 12, 2024. [Online]. Available: <https://arxiv.org/abs/2307.03172>
- **[Liu et al., ProRL]** M. Liu, S. Diao, X. Lu, *et al.*, "ProRL: Prolonged reinforcement learning expands reasoning boundaries in large language models," in *Proc. NeurIPS*, 2025. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2025/hash/1a22b912945fb7c0bdd079e792b31b6f-Abstract-Conference.html
- **[Liu et al., ScaleCUA]** Z. Liu, *et al.*, "ScaleCUA: Scaling open-source computer use agents with cross-platform data," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2509.15221>
- **[Liu et al., Who&When Pro]** J. Liu, H. Xi, S. Zhang, *et al.*, "Who&When Pro: Can LLMs really attribute failures in AI agents?," in *Proc. ICML Workshop on Failure Modes of Agentic AI*, 2026. [Online]. Available: <https://arxiv.org/abs/2607.09996>
- **[Lu et al., MathVista]** P. Lu, H. Bansal, *et al.*, "MathVista: Evaluating mathematical reasoning of foundation models in visual contexts," in *Proc. ICLR*, 2024. [Online]. Available: <https://openreview.net/forum?id=KUNzEQMWU7>
- **[Lu et al., WebGen-Bench]** Z. Lu, *et al.*, "WebGen-Bench: Evaluating LLMs on generating interactive and functional websites from scratch," in *Proc. NeurIPS Datasets and Benchmarks Track*, 2025. [Online]. Available: <https://arxiv.org/abs/2505.03733>

- **[Lu et al., WebLINX]** X. H. Lù, Z. Kasner, and S. Reddy, "WebLINX: Real-world website navigation with multi-turn dialogue," in *Proc. ICML*, 2024. [Online]. Available: <https://arxiv.org/abs/2402.05930>
- **[Luo et al., Agentix]** M. Luo, X. Shi, C. Cai, *et al.*, "Agentix: An efficient serving engine for LLM agents as general programs," in *Proc. USENIX NSDI*, 2026, pp. 2443–2459. [Online]. Available: <https://www.usenix.org/conference/nsdi26/presentation/luo>
- **[Lupinacci et al., dark side]** M. Lupinacci, F. A. Pironti, F. Blefari, *et al.*, "The dark side of LLMs: Agent-based attack vectors for system-level compromise," in *Lecture Notes in Computer Science*, Springer, 2026. [Online]. Available: <https://arxiv.org/abs/2507.06850>
- **[Lynch et al., Agentic misalignment]** A. Lynch, *et al.*, "Agentic misalignment: How LLMs could be insider threats," *arXiv:2510.05179*, Oct. 2025. [Online]. Available: <https://arxiv.org/abs/2510.05179>
- **[Ma et al., AgentBoard]** C. Ma, *et al.*, "AgentBoard: An analytical evaluation board of multi-turn LLM agents," in *Proc. NeurIPS Datasets and Benchmarks Track*, 2024. [Online]. Available: <https://arxiv.org/abs/2401.13178>
- **[Ma et al., DoVer]** M. Ma, J. Zhang, F. Yang, Y. Kang, Q. Lin, S. Rajmohan, and D. Zhang, "DoVer: Intervention-driven auto debugging for LLM multi-agent systems," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2512.06749>
- **[MacDiarmid et al., Emergent Misalignment]** M. MacDiarmid, B. Wright, J. Uesato, *et al.* (Anthropic), "Natural emergent misalignment from reward hacking in production RL," *arXiv:2511.18397*, Nov. 2025. [Online]. Available: <https://arxiv.org/abs/2511.18397>
- **[Madaan et al., Self-Refine]** A. Madaan, N. Tandon, P. Gupta, *et al.*, "Self-Refine: Iterative refinement with self-feedback," in *Proc. NeurIPS*, 2023. [Online]. Available: <https://arxiv.org/abs/2303.17651>
- **[Magar & Schwartz, Data contamination]** I. Magar and R. Schwartz, "Data contamination: From memorization to exploitation," in *Proc. ACL*, 2022. [Online]. Available: <https://arxiv.org/abs/2203.08242>
- **[Maharana et al., LoCoMo]** A. Maharana, D.-H. Lee, S. Tulyakov, M. Bansal, F. Barbieri, and Y. Fang, "Evaluating very long-term conversational memory of LLM agents (LoCoMo)," in *Proc. ACL*, 2024. [Online]. Available: <https://aclanthology.org/2024.acl-long.747/>

- **[Managed agent runtimes]** Claude Managed Agents, Amazon Bedrock AgentCore, Google Gemini Enterprise Agent Platform: managed agent runtimes whose tracing is a deployment setting, referenced in *Chapter 12*.
- **[Manus, Context Engineering]** Manus (Y. Ji), "Context engineering for AI agents: Lessons from building Manus," *Manus Blog*, Jul. 2025. [Online]. Available: <https://manus.im/blog/Context-Engineering-for-AI-Agents-Lessons-from-Building-Manus>
- **[Mao & Mirhoseini, DeLM]** Y. Mao and A. Mirhoseini, "Decentralized multi-agent systems with shared context (DeLM)," in *Proc. COLM Workshop on Lifelong Agents*, 2026. [Online]. Available: <https://arxiv.org/abs/2606.10662>
- **[Marquet, Turn the Ship]** L. D. Marquet, *Turn the Ship Around! A True Story of Turning Followers into Leaders*. New York: Portfolio, 2012. [Online]. Available: <https://www.penguinrandomhouse.com/books/314163/turn-the-ship-around-by-l-david-marquet/>
- **[Martin et al., Scaling Managed Agents]** L. Martin, G. Cemaj, and M. Cohen (Anthropic), "Scaling Managed Agents: Decoupling the brain from the hands," *Anthropic Engineering*, Apr. 2026. [Online]. Available: <https://www.anthropic.com/engineering/managed-agents>
- **[Martin, Design principles]** R. C. Martin, "Design principles and design patterns (including the Interface Segregation Principle)," *Object Mentor*, 2000. [Online]. Available: https://staff.cs.utu.fi/~jounsmcd/doos_06/material/DesignPrinciplesAndPatterns.pdf
- **[Mastercard, Verifiable Intent]** Mastercard (maintainer), "Verifiable Intent: Open specification for cryptographic agent authorization in commerce," Draft v0.1, *GitHub*, Feb. 2026. [Online]. Available: <https://github.com/agent-intent/verifiable-intent>
- **[Masters, Rasch model]** G. N. Masters, "A Rasch model for partial credit scoring," *Psychometrika*, vol. 47, no. 2, pp. 149–174, Jun. 1982. [Online]. Available: <https://doi.org/10.1007/BF02296272>
- **[Mazeika et al., Remote Labor Index]** M. Mazeika, *et al.* (Center for AI Safety and Scale AI), "Remote Labor Index: Measuring AI automation of remote work," *Center for AI Safety and Scale AI report*, Oct. 2025. [Online]. Available: <https://arxiv.org/abs/2510.26787>

- **[McClelland et al., Complementary Learning]** J. L. McClelland, B. L. McNaughton, and R. C. O'Reilly, "Why there are complementary learning systems in the hippocampus and neocortex: Insights from the successes and failures of connectionist models of learning and memory," *Psychological Review*, vol. 102, no. 3, pp. 419–457, Jul. 1995. [Online]. Available: <https://doi.org/10.1037/0033-295X.102.3.419>
- **[McKinsey, Agentic Advantage]** McKinsey & Company (QuantumBlack), "Seizing the agentic AI advantage," *McKinsey Insights*, 2025. [Online]. Available: <https://www.mckinsey.com/capabilities/quantumblack/our-insights/seizing-the-agentic-ai-advantage>
- **[MCP Spec 2025-03-26]** ModelContextProtocol Authors, "MCP specification revision 2025-03-26: tool annotations (readOnlyHint, destructiveHint, idempotentHint, openWorldHint) and the requirement that clients treat them as untrusted unless the server is trusted," *modelcontextprotocol.io*, Mar. 2025. [Online]. Available: <https://modelcontextprotocol.io/specification/2025-03-26/server/tools>
- **[MCP Spec 2026-07-28]** ModelContextProtocol Authors, "MCP specification revision 2026-07-28: stateless core (SEP-2567 / SEP-2575), official tasks extension (SEP-2663), trace-context conventions (SEP-414), OAuth 2.1 / RFC 8707 authorization, and formal feature lifecycle with a twelve-month deprecation window (SEP-2596)," *modelcontextprotocol.io*, Jul. 2026. [Online]. Available: <https://modelcontextprotocol.io/specification/2026-07-28/changelog>
- **[MCP-Bench, MCP-Bench]** Z. Wang, *et al.*, "MCP-Bench: Benchmarking tool-using LLM agents with complex real-world tasks via MCP servers," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2508.20453>
- **[Meincke et al., Prompting Science 2]** L. Meincke, E. R. Mollick, L. Mollick, and D. Shapiro, "Prompting science report 2: The decreasing value of chain of thought in prompting," *Wharton Generative AI Labs technical report*, Jun. 2025. [Online]. Available: <https://arxiv.org/abs/2506.07142>
- **[Mem0, Mem0]** P. Chhikara, D. Khant, S. Aryan, T. Singh, and D. Yadav, "Mem0: Building production-ready AI agents with scalable long-term memory," in *Proc. ECAI*, 2025. [Online]. Available: <https://arxiv.org/abs/2504.19413>
- **[Menlo Ventures, State of Generative AI 2025]** Menlo Ventures, "2025: The state of generative AI in the enterprise," *Menlo Ventures Perspective*, Dec. 2025. [Online]. Available: <https://menlovc.com/perspective/2025-the-state-of-generative-ai-in-the-enterprise/>

- **[Mercor, Enterprise Evals]** Mercor, "Enterprise evals," *Mercor*, continuously updated; checked Sep. 2026. [Online]. Available: <https://www.mercor.com/enterprise-evals/>
- **[Meta, Llama Guard 3]** Meta AI, "Llama Guard 3 / Llama Guard 4 / Prompt Guard: Open-source safety classifiers," *Meta AI Releases*, 2024–2025. [Online]. Available: <https://github.com/meta-llama/PurpleLlama>
- **[Meta, LlamaFirewall]** Meta AI, "LlamaFirewall: An open-source guardrail system for AI agents (PromptGuard 2, AlignmentCheck, CodeShield)," *PurpleLlama documentation*, 2025. [Online]. Available: <https://meta-llama.github.io/PurpleLlama/LlamaFirewall/docs/documentation/about-llamafirewall>
- **[METR, Horizon Assumptions]** A. Barry (METR), "Research note: Impact of modeling assumptions on time horizon results," *METR*, Mar. 2026. [Online]. Available: <https://metr.org/notes/2026-03-20-impact-of-modelling-assumptions-on-time-horizon-results/>
- **[METR, Reward hacking]** METR, "Recent frontier models are reward hacking," *METR Blog*, Jun. 2025. [Online]. Available: <https://metr.org/blog/2025-06-05-recent-reward-hacking/>
- **[METR, Time Horizon 1.1]** METR, "Time Horizon 1.1," *METR Blog*, Jan. 2026. [Online]. Available: <https://metr.org/blog/2026-1-29-time-horizon-1-1/>
- **[METR, Time Horizons]** METR, "Task-completion time horizons of frontier AI models," *METR*, continuously updated. [Online]. Available: <https://metr.org/time-horizons/>
- **[METR, Uplift Update]** METR, "We are changing our developer productivity experiment design," *METR Blog*, Feb. 2026. [Online]. Available: <https://metr.org/blog/2026-02-24-uplift-update/>
- **[Mialon et al., GAIA]** G. Mialon, C. Fourrier, C. Swift, T. Wolf, Y. LeCun, and T. Scialom, "GAIA: A benchmark for general AI assistants," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2311.12983>
- **[Microsoft Agent Framework, Microsoft Agent Framework]** Microsoft, "Microsoft Agent Framework — agent-first stack consolidating AutoGen and Semantic Kernel," *Microsoft Learn / GitHub*, continuously updated. [Online]. Available: <https://github.com/microsoft/agent-framework>
- **[Microsoft Durable Functions, Durable Functions]** Microsoft, "Durable Functions — stateful workflows in serverless Azure Functions," *Microsoft Learn*, continuously

updated. [Online]. Available: <https://learn.microsoft.com/en-us/azure/durable-task/durable-functions/durable-functions-overview>

- **[Microsoft Presidio, Presidio]** Data Privacy Stack (formerly Microsoft), "Presidio: Open-source PII detection and anonymization framework," *presidio.dataprivacystack.org*, continuously updated. [Online]. Available: <https://presidio.dataprivacystack.org/>
- **[Microsoft Semantic Kernel, Semantic Kernel]** Microsoft, "Semantic Kernel — model-agnostic SDK for composing prompts, tools, and planners as skills," *Microsoft Learn*, continuously updated. [Online]. Available: <https://learn.microsoft.com/en-us/semantic-kernel/>
- **[Microsoft, Agent Framework Overview]** Microsoft, "Microsoft Agent Framework overview: the direct successor to Semantic Kernel and AutoGen, created by the same teams, with migration guides from both," *Microsoft Learn*, continuously updated. [Online]. Available: <https://learn.microsoft.com/en-us/agent-framework/overview/>
- **[Microsoft, Agent Network Red Team]** G. Bansal, S. Mirza, K. Hines, *et al.* (Microsoft Research), "Red-teaming a network of agents: Understanding what breaks when AI agents interact at scale," *Microsoft Research Blog*, Apr. 2026. [Online]. Available: <https://www.microsoft.com/en-us/research/blog/red-teaming-a-network-of-agents-understanding-what-breaks-when-ai-agents-interact-at-scale/>
- **[Microsoft, AI Red Team]** B. Bullwinkel and R. Shankar Siva Kumar (Microsoft), "3 takeaways from red teaming 100 generative AI products," *Microsoft Security Blog*, Jan. 2025. [Online]. Available: <https://www.microsoft.com/en-us/security/blog/2025/01/13/3-takeaways-from-red-teaming-100-generative-ai-products/>
- **[Microsoft, AI Red Teaming Agent]** Microsoft, "AI Red Teaming Agent: PyRIT attack strategies run against Foundry-hosted agents and scored by attack success rate," *Microsoft Learn*, continuously updated. [Online]. Available: <https://learn.microsoft.com/en-us/azure/foundry/concepts/ai-red-teaming-agent>
- **[Microsoft, AutoGen]** Microsoft, "AutoGen AgentChat: stateful agents and managed teams built on AutoGen Core," *AutoGen Documentation*, continuously updated. [Online]. Available: <https://microsoft.github.io/autogen/stable/user-guide/agentchat-user-guide/>
- **[Microsoft, Content Safety]** Microsoft, "Azure AI Content Safety: Prompt Shields and groundedness detection (managed runtime checks, with a low-latency non-reasoning

mode and an optional groundedness correction)," *Microsoft Learn*, continuously updated. [Online]. Available: <https://learn.microsoft.com/en-us/azure/ai-services/content-safety/>

- **[Microsoft, CORPGEN]** A. Jaye, N. B. Kumankumah, C. Biringa, *et al.* (Microsoft), "CORPGEN advances AI agents for real work," *Microsoft Research Blog*, Feb. 2026. [Online]. Available: <https://www.microsoft.com/en-us/research/blog/corpgen-advances-ai-agents-for-real-work/>
- **[Microsoft, Entra Agent ID]** Microsoft, "Microsoft Entra Agent ID: agent identities, OAuth protocols, and delegation for AI agents," *Microsoft Learn*, 2026. [Online]. Available: <https://learn.microsoft.com/en-us/entra/agent-id/agent-identities>
- **[Microsoft, Magentic Orchestration]** Microsoft, "Magentic orchestration — manager-led planning with optional human plan review (RequirePlanSignoff / enable_plan_review)," *Microsoft Learn*, continuously updated. [Online]. Available: <https://learn.microsoft.com/en-us/agent-framework/workflows/orchestrations/magentic>
- **[Microsoft, Ten Months in dotnet/runtime]** S. Toub, "Ten months with Copilot coding agent in dotnet/runtime," *.NET Blog*, Microsoft, Mar. 2026. [Online]. Available: <https://devblogs.microsoft.com/dotnet/ten-months-with-cca-in-dotnet-runtime/>
- **[Miller, Error Bars]** E. Miller, "Adding error bars to evals: A statistical approach to language model evaluations," *arXiv:2411.00640*, Nov. 2024. [Online]. Available: <https://arxiv.org/abs/2411.00640>
- **[Mireshghallah et al., ConfAide]** N. Mireshghallah, H. Kim, X. Zhou, Y. Tsvetkov, M. Sap, R. Shokri, and Y. Choi, "Can LLMs keep a secret? Testing privacy implications of language models via contextual integrity theory (ConfAide)," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2310.17884>
- **[MITRE ATLAS, ATLAS]** MITRE, "Adversarial Threat Landscape for Artificial-Intelligence Systems (ATLAS) knowledge base," MITRE, continuously updated. [Online]. Available: <https://atlas.mitre.org/>
- **[Modal Docs, Modal]** Modal Labs, "Modal — serverless GPU and CPU containers for AI workloads," *Modal Documentation*, continuously updated. [Online]. Available: <https://modal.com/docs>

- **[Moran, your multi-agent system]** M. Moran, "Why your multi-agent system is failing: Escaping the 17× error trap of the 'bag of agents'," *Towards Data Science*, 2026. [Online]. Available: <https://towardsdatascience.com/why-your-multi-agent-system-is-failing-escaping-the-17x-error-trap-of-the-bag-of-agents/>
- **[Motwani et al., Secret collusion]** S. Motwani, *et al.*, "Secret collusion among AI agents: Multi-agent deception via steganography," in *Proc. NeurIPS*, 2024. [Online]. Available: <https://arxiv.org/abs/2402.07510>
- **[Mozannar et al., Magentic-UI]** H. Mozannar, G. Bansal, C. Tan, *et al.* (Microsoft Research), "Magentic-UI: Towards human-in-the-loop agentic systems," *Microsoft Research Technical Report MSR-TR-2025-40*, Jul. 2025. [Online]. Available: <https://www.microsoft.com/en-us/research/publication/magentic-ui-report/>
- **[Nainani et al., Timing Analysis]** J. Nainani, C.-T. Ho, A. Dhurka, and H. Ren, "Timing analysis agent: Autonomous multi-corner multi-mode (MCM) timing debugging with timing debug relation graph," *arXiv:2504.11502*, Apr. 2025. [Online]. Available: <https://arxiv.org/abs/2504.11502>
- **[Nasdaq Verafin, Agentic Workforce]** Nasdaq Verafin, "Nasdaq Verafin announces expansion of its agentic AI workforce," *Verafin news*, Jun. 2026. [Online]. Available: <https://verafin.com/news/nasdaq-verafin-announces-expansion-of-its-agentic-ai-workforce/>
- **[Nasr et al., Attacker Moves Second]** M. Nasr, N. Carlini, C. Sitawarin, *et al.*, "The attacker moves second: Stronger adaptive attacks bypass defenses against LLM jailbreaks and prompt injections," *arXiv:2510.09023*, Oct. 2025. [Online]. Available: <https://arxiv.org/abs/2510.09023>
- **[Nath and Krishnaswamy, Partner-Aware]** A. Nath and N. Krishnaswamy, "Learning partner-aware collaborators in multi-party collaboration," in *Proc. NeurIPS*, 2025. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2025/hash/be3e865f449b78ad35d711ac8e418c6d-Abstract-Conference.html
- **[NATS JetStream, NATS JetStream]** Synadia / NATS Authors, "NATS JetStream — persistence layer for the NATS messaging system with at-least-once / exactly-once delivery semantics," *nats.io*, continuously updated. [Online]. Available: <https://docs.nats.io/nats-concepts/jetstream>
- **[Ng, blog series]** A. Ng / DeepLearning.AI, "AI agentic design patterns (blog series)," *The Batch*, Apr. 2024. [Online]. Available: <https://www.deeplearning.ai/the-batch/how-agents-can-improve-llm-performance/>

- **[Ni et al., Collaborative Reasoner]** A. Ni, R. Desai, Y. Li, *et al.*, "Collaborative Reasoner: Self-improving social agents with synthetic conversations," in *Proc. NeurIPS*, 2025. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2025/hash/221ae0f5de12f7b9803af2656ee7902d-Abstract-Conference.html
- **[Ning et al., MC-Search]** X. Ning, D. Fu, T. Wei, *et al.*, "MC-Search: Evaluating and enhancing multimodal agentic search with structured long reasoning chains," in *Proc. ICLR Oral*, 2026. [Online]. Available: <https://openreview.net/forum?id=JEGDp1E40H>
- **[Nissenbaum, Contextual Integrity]** H. Nissenbaum, "Privacy as contextual integrity," *Washington Law Review*, vol. 79, no. 1, pp. 119–157, Feb. 2004. [Online]. Available: <https://digitalcommons.law.uw.edu/wlr/vol79/iss1/10/>
- **[Novikov et al., AlphaEvolve]** A. Novikov, N. Vů, M. Eisenberger, *et al.*, "AlphaEvolve: A coding agent for scientific and algorithmic discovery," *arXiv:2506.13131*, Jun. 2025. [Online]. Available: <https://arxiv.org/abs/2506.13131>
- **[NVIDIA garak, probes, generators, detectors]** NVIDIA, "garak: generative AI red-teaming and assessment kit — open-source LLM vulnerability scanner (probes, generators, detectors)," *GitHub*, continuously updated. [Online]. Available: <https://github.com/NVIDIA/garak>
- **[NVIDIA TensorRT-LLM, TensorRT-LLM]** NVIDIA, "TensorRT-LLM — a toolbox for optimized large language model inference on NVIDIA GPUs," *NVIDIA Developer Documentation*, continuously updated. [Online]. Available: <https://nvidia.github.io/TensorRT-LLM/>
- **[NVIDIA, DCGM]** NVIDIA, "NVIDIA Data Center GPU Manager (DCGM) user guide: feature overview and profiling metrics (DCGM_FI_DEV_GPU_UTIL, DCGM_FI_PROF_SM_ACTIVE, DCGM_FI_PROF_PIPE_TENSOR_ACTIVE)," *NVIDIA Data Center Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://docs.nvidia.com/datacenter/dcgm/latest/user-guide/feature-overview.html>
- **[NVIDIA, NeMo Guardrails]** NVIDIA, "NeMo Guardrails and Colang: Programmable LLM safety rails," *NVIDIA Open-Source Documentation*, 2024–2025. [Online]. Available: <https://github.com/NVIDIA-NeMo/Guardrails>
- **[NVlabs, VerilogCoder source]** NVIDIA Research, "VerilogCoder source code," *GitHub*, commit 8b13108, checked Sep. 2026. [Online]. Available: <https://github.com/NVlabs/VerilogCoder/tree/8b13108869f276c7b644dd88beabcb401a5cfa92>

- **[Ogenrwot & Businge, AgenticFlict]** D. Ogenrwot and J. Businge, "AgenticFlict: A large-scale dataset of merge conflicts in AI coding agent pull requests on GitHub," in *Proc. ACM Int. Conf. AI-Powered Software (AIware)*, Montreal, QC, Canada, Jul. 2026, doi: 10.1145/3805760.3814923. [Online]. Available: <https://arxiv.org/abs/2604.03551>
- **[Ogenrwot & Businge, Modify Code]** D. Ogenrwot and J. Businge, "How AI coding agents modify code: A large-scale study of GitHub pull requests," in *Proc. 23rd Int. Conf. Mining Software Repositories (MSR), Mining Challenge Track*, Apr. 2026. [Online]. Available: <https://arxiv.org/abs/2601.17581>
- **[Ollama, Ollama]** Ollama, "Ollama — run large language models locally," *ollama.com*, continuously updated. [Online]. Available: <https://ollama.com/>
- **[OneUptime, GDPR-compliant OpenTelemetry pipeline]** OneUptime, "Building a GDPR-compliant OpenTelemetry pipeline," *OneUptime Blog*, Feb. 2026. [Online]. Available: <https://oneuptime.com/blog/post/2026-02-06-opentelemetry-pipeline-gdpr-compliant/view>
- **[Ong et al., RouteLLM]** I. Ong, A. Almahairi, V. Wu, *et al.*, "RouteLLM: Learning to route LLMs from preference data," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2406.18665>
- **[OpenAI, Agents SDK]** OpenAI, "OpenAI Agents SDK: Handoffs, guardrails, and tracing," *GitHub*, 2025. [Online]. Available: <https://github.com/openai/openai-agents-python>
- **[OpenAI, BrowseComp]** OpenAI, "BrowseComp: A benchmark for browsing agents," *OpenAI*, Apr. 2025. [Online]. Available: <https://openai.com/index/browsecomp/>
- **[OpenAI, Codex Approvals]** OpenAI, "Agent approvals & security: sandbox modes, approval policies, and protected paths," *Codex documentation*, read 21 Sep. 2026. [Online]. Available: <https://learn.chatgpt.com/docs/agent-approvals-security>
- **[OpenAI, Codex Auto-review]** OpenAI, "Auto-review," *Codex documentation*, read 17 Sep. 2026. [Online]. Available: <https://learn.chatgpt.com/docs/sandboxing/auto-review>
- **[OpenAI, Codex Harness]** OpenAI, "Codex: the open agent harness and terminal coding agent," *GitHub*, 2026. [Online]. Available: <https://github.com/openai/codex>
- **[OpenAI, Codex Sandbox Advisory]** OpenAI, "Sandbox bypass due to bug in path configuration logic," *GitHub security advisory GHSA-w5fx-fh39-j5rw (CVE-2025-59532)*, Sep. 2025. [Online]. Available: <https://github.com/openai/codex/security/advisories/GHSA-w5fx-fh39-j5rw>

- **[OpenAI, Codex-Max System Card]** OpenAI, "GPT-5.1-Codex-Max system card," *OpenAI*, Nov. 2025. [Online]. Available: https://cdn.openai.com/pdf/2a7d98b1-57e5-4147-8d0e-683894d782ae/5p1_codex_max_card_03.pdf
- **[OpenAI, Delegated Payment]** OpenAI and Stripe, "Delegated payment spec," *Agentic Commerce Protocol documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://developers.openai.com/commerce/specs/payment>
- **[OpenAI, Deprecations]** OpenAI, "Deprecations — Agent Builder wind-down (announced Jun. 2026, shutdown Nov. 2026; migration to the Agents SDK or ChatGPT Workspace Agents) and Assistants API sunset (Aug. 2026, superseded by the Responses and Conversations APIs)," *OpenAI Developer Documentation*, continuously updated. [Online]. Available: <https://developers.openai.com/api/docs/deprecations>
- **[OpenAI, Flex Processing]** OpenAI, "Flex processing — service tiers at Batch API rates for lower-priority workloads," *OpenAI Platform Documentation*, continuously updated. [Online]. Available: <https://developers.openai.com/api/docs/guides/flex-processing>
- **[OpenAI, Guardian Delegation Review]** OpenAI, "Preserve root user authorization in subagent Guardian reviews," *openai/codex*, pull request 39975, Aug. 2026. [Online]. Available: <https://github.com/openai/codex/pull/39975>
- **[OpenAI, Guardian Sender Context]** OpenAI, "Include sender user messages in Guardian delegation reviews," *openai/codex*, pull request 46179, Sep. 2026. [Online]. Available: <https://github.com/openai/codex/pull/46179>
- **[OpenAI, Hugging Face incident]** OpenAI, "OpenAI – Hugging Face incident: technical report," *OpenAI*, Aug. 2026. [Online]. Available: <https://cdn.openai.com/pdf/67869394-cb91-4c12-888c-5cbd85c7814c/OpenAI-Hugging-Face%20Incident-Technical-Report.pdf>
- **[OpenAI, Human in the Loop]** OpenAI, "Human in the loop — pausing a run on a tool approval and resuming it from a serialized RunState," *OpenAI Agents SDK Documentation*, continuously updated. [Online]. Available: <https://openai.github.io/openai-agents-js/guides/human-in-the-loop/>
- **[OpenAI, Open Responses]** OpenAI, Nvidia, Vercel, OpenRouter, Hugging Face, LM Studio, Ollama, and vLLM, "Open Responses: a multi-provider specification on top of the OpenAI Responses API," *OpenAI Developers Changelog*, 2026. [Online]. Available: <https://developers.openai.com/api/docs/changelog>

- **[OpenAI, Orchestrating Agents]** I. Bigio (OpenAI), "Orchestrating agents: Routines and handoffs," *OpenAI Cookbook*, Oct. 2024. [Online]. Available: https://developers.openai.com/cookbook/examples/orchestrating_agents
- **[OpenAI, Priority Processing]** OpenAI, "Priority processing (Fast mode) — consistent low latency at a premium per-token rate," *OpenAI Platform Documentation*, continuously updated. [Online]. Available: <https://developers.openai.com/api/docs/guides/fast-mode>
- **[OpenAI, Prompt Caching]** OpenAI, "Prompt caching," *OpenAI Platform Documentation*, Oct. 2024. [Online]. Available: <https://developers.openai.com/api/docs/guides/prompt-caching>
- **[OpenAI, Reasoning]** OpenAI, "Reasoning models: reasoning tokens, context, and billing," *OpenAI API Documentation*, continuously updated. [Online]. Available: <https://developers.openai.com/api/docs/guides/reasoning>
- **[OpenAI, SWE-bench Retired]** OpenAI, "Why SWE-bench Verified no longer measures frontier coding capabilities," *OpenAI*, Feb. 2026. [Online]. Available: <https://openai.com/index/why-we-no-longer-evaluate-swe-bench-verified/>
- **[OpenAI, SWE-bench Verified]** OpenAI, "Introducing SWE-bench Verified," *OpenAI*, Aug. 2024. [Online]. Available: <https://openai.com/index/introducing-swe-bench-verified/>
- **[OpenClaw, Config Mutation Advisory]** OpenClaw Contributors, "Gateway config mutation guard allowed unsafe model-driven config writes," *GitHub security advisory GHSA-cwj3-vqpp-pmxr*, Apr. 2026. [Online]. Available: <https://github.com/openclaw/openclaw/security/advisories/GHSA-cwj3-vqpp-pmxr>
- **[OpenClaw, Heartbeat Advisory]** OpenClaw Contributors, "Heartbeat context inheritance bypasses sandbox via senderIsOwner escalation," *GitHub security advisory GHSA-g5cg-8x5w-7jpm*, Mar. 2026. [Online]. Available: <https://github.com/openclaw/openclaw/security/advisories/GHSA-g5cg-8x5w-7jpm>
- **[OpenClaw, OpenClaw]** OpenClaw Contributors, "OpenClaw: your own personal AI assistant, any OS, any platform," *GitHub*, 2026. [Online]. Available: <https://github.com/openclaw/openclaw>
- **[OpenClaw, VirusTotal Partnership]** P. Steinberger, J. O'Reilly, and B. Quintero, "OpenClaw partners with VirusTotal for skill security," *OpenClaw Blog*, Feb. 2026. [Online]. Available: <https://openclaw.ai/blog/virustotal-partnership>

- **[OpenCode, OpenCode]** OpenCode Contributors, "OpenCode: the open source AI coding agent," *GitHub*, commit 82d4c89, checked Sep. 2026. [Online]. Available: <https://github.com/anomalyco/opencode/tree/82d4c89031ed30da1cf22ba62419fec520df8d27>
- **[OpenCode, Server Advisory]** OpenCode Contributors, "Unauthenticated HTTP server allows arbitrary command execution," *GitHub security advisory GHSA-vxw4-wv6m-9hhh (CVE-2026-22812)*, Jan. 2026. [Online]. Available: <https://github.com/anomalyco/opencode/security/advisories/GHSA-vxw4-wv6m-9hhh>
- **[OpenCost, Specification]** OpenCost Authors (CNCF incubating project), "OpenCost specification: Kubernetes cost monitoring and allocation, including cluster idle cost," *OpenCost Specification*, continuously updated. [Online]. Available: <https://opencost.io/docs/specification>
- **[OpenID Foundation, Agentic Identity]** T. South (lead ed.), *et al.*, "Identity management for agentic AI: The new frontier of authorization, authentication, and security for an AI agent world," *OpenID Foundation*, Oct. 2025. [Online]. Available: <https://openid.net/wp-content/uploads/2025/10/Identity-Management-for-Agentic-AI.pdf>
- **[OpenSSF, Model Signing]** M. Maruseac (Google Open Source Security Team), "Taming the wild west of ML: Practical model signing with Sigstore," *Sigstore Blog*, Apr. 2025. [Online]. Available: <https://blog.sigstore.dev/model-transparency-v1.0/>
- **[OpenTelemetry]** OpenTelemetry Authors, "OpenTelemetry specification," *opentelemetry.io*, continuously updated. [Online]. Available: <https://opentelemetry.io/docs/specs/otel/>
- **[OpenTelemetry GenAI]** OpenTelemetry Authors, "Semantic conventions for generative AI (gen_ai.* attributes; moved to the dedicated GenAI semantic-conventions repository with semantic-conventions v1.42.0, June 2026)," *GitHub*, 2026. [Online]. Available: <https://github.com/open-telemetry/semantic-conventions-genai>
- **[OpenTelemetry MCP]** OpenTelemetry Authors, "MCP span and attribute conventions (Development stability, merged Jan 2026; moved to the dedicated GenAI semantic-conventions repository June 2026)," *GitHub*, 2026. [Online]. Available: <https://github.com/open-telemetry/semantic-conventions-genai/tree/main/docs/gen-ai>

- **[OpenTelemetry, Database Spans]** OpenTelemetry Authors, "Semantic conventions for database client spans (db.* attributes; Stable)," *OpenTelemetry Specification*, continuously updated; checked Sep. 2026. [Online]. Available: <https://opentelemetry.io/docs/specs/semconv/db/database-spans/>
- **[OpenTelemetry, GenAI Cost Proposal]** OpenTelemetry Authors, "Add per-operation cost conventions (gen_ai.usage.cost.*)," pull request #443, *open-telemetry/semantic-conventions-genai*, GitHub, Aug. 2026 (open, not merged, when read in September 2026). [Online]. Available: <https://github.com/open-telemetry/semantic-conventions-genai/pull/443>
- **[OpenTelemetry, Hardware Metrics]** OpenTelemetry Authors, "Semantic conventions for GPU metrics (hw.gpu.*; Development stability)," *OpenTelemetry Specification*, continuously updated; checked Sep. 2026. [Online]. Available: <https://opentelemetry.io/docs/specs/semconv/hardware/gpu/>
- **[OpenTelemetry, Span Event Deprecation]** L. Molkova, R. Pająk, and T. Stalnakier, "Deprecating Span Events API," *OpenTelemetry Blog*, Mar. 17, 2026. [Online]. Available: <https://opentelemetry.io/blog/2026/deprecating-span-events/>
- **[OpenTelemetry, Tail Sampling]** OpenTelemetry Authors, "Tail sampling processor," *OpenTelemetry Collector Contrib (GitHub)*, continuously updated; checked Sep. 2026. [Online]. Available: <https://github.com/open-telemetry/opentelemetry-collector-contrib/blob/main/processor/tailsamplingprocessor/README.md>
- **[Orlanski et al., LibraryDesignBench]** G. Orlanski, A. L. Zhang, A. Trost, V. S. Chen, F. Sala, A. Albarghouthi, and L. Schmidt, "Can agents design libraries for agents?," *arXiv: 2609.36730*, Sep. 2026. [Online]. Available: <https://arxiv.org/abs/2609.36730>
- **[Orlanski et al., SlopCodeBench]** G. Orlanski, D. Roy, A. Yun, *et al.*, "SlopCodeBench: Benchmarking how coding agents degrade over long-horizon iterative tasks," *arXiv: 2603.24755*, May 2026. [Online]. Available: <https://arxiv.org/abs/2603.24755>
- **[Outshift, Hands-on with CAIPE]** Outshift by Cisco, "Hands-on with CAIPE: Building an open source, multi-agent system for platform engineering," *Outshift Blog*, 2025. [Online]. Available: <https://outshift.cisco.com/blog/inside-outshift/caipe-building-open-source-multi-agent-systems-for-platform-engineering>
- **[Outshift, Splunk Forge]** S. Aradhyula, B. Hartwell, H. Kalpage, and K. Kantesaria, "How Splunk Cloud Platform team used CAIPE to power its always-on internal AI assistant," *Outshift by Cisco Blog*, Apr. 2026, updated Jun. 2026. [Online]. Available: <https://outshift.cisco.com/blog/ai-ml/how-splunk-built-forge-on-caipe>

- **[Ouyang et al., InstructGPT]** L. Ouyang, J. Wu, X. Jiang, *et al.*, "Training language models to follow instructions with human feedback (InstructGPT)," in *Proc. NeurIPS*, 2022. [Online]. Available: <https://arxiv.org/abs/2203.02155>
- **[Ouyang et al., ReasoningBank]** S. Ouyang, *et al.*, "ReasoningBank: Scaling agent self-evolving with reasoning memory," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2509.25140>
- **[OWASP, Agentic Top 10]** OWASP Foundation, "OWASP Top 10 for Agentic Applications for 2026," Dec. 2025. [Online]. Available: <https://genai.owasp.org/resource/owasp-top-10-for-agentic-applications-for-2026/>
- **[OWASP, LLM Top 10]** OWASP Foundation, "OWASP Top 10 for LLM Applications 2025 (LLM01 Prompt Injection through LLM10 Unbounded Consumption)," Nov. 2024. [Online]. Available: <https://genai.owasp.org/llm-top-10/>
- **[OWASP, LLM Top 10 2026]** OWASP Foundation, "OWASP Top 10 for LLM Applications 2026," *OWASP Gen AI Security Project*, Aug. 2026. [Online]. Available: <https://genai.owasp.org/resource/owasp-genai-llm-top-10-2026/>
- **[OWASP, NHI]** OWASP Foundation, "OWASP Non-Human Identities (NHI) Top 10 (2025): NHI1 Improper Offboarding through NHI10 Human Use of NHI," 2025. [Online]. Available: <https://owasp.org/projects/non-human-identities-top-10>
- **[OWASP, Unbounded Consumption]** OWASP Foundation, "OWASP Top 10 for LLM Applications 2025 — LLM10:2025 Unbounded Consumption," 2025. [Online]. Available: <https://genai.owasp.org/llmrisk/llm102025-unbounded-consumption/>
- **[Packer et al., MemGPT]** C. Packer, *et al.*, "MemGPT: Towards LLMs as operating systems," *arXiv:2310.08560*, Oct. 2023. [Online]. Available: <https://arxiv.org/abs/2310.08560>
- **[Pan et al., KVFlow]** Z. Pan, *et al.*, "KVFlow: Efficient prefix caching for accelerating LLM-based multi-agent workflows," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2507.07400>
- **[Pandey, Internet of Agents]** V. Pandey, "The Internet of Agents: an open, interoperable internet for agent-agent and agent-human quantum-safe communication," v4, *Outshift by Cisco White Paper*, 2025. [Online]. Available: <https://outshift.cisco.com/the-internet-of-agents>
- **[Pang et al., Interactive Reasoning]** R. Y. Pang, K. J. K. Feng, S. Feng, *et al.*, "Interactive reasoning: Visualizing and controlling chain-of-thought reasoning in large language

- models," in *Proc. ACM IUI*, pp. 852–867, 2026. [Online]. Available: <https://doi.org/10.1145/3742413.3789091>
- **[Panickssery et al., LLM evaluators recognize]** A. Panickssery, S. R. Bowman, and S. Feng, "LLM evaluators recognize and favor their own generations," in *Proc. NeurIPS*, 2024. [Online]. Available: <https://arxiv.org/abs/2404.13076>
 - **[Paradis et al., AI development speed]** E. Paradis, *et al.*, "How much does AI impact development speed? An enterprise-based randomized controlled trial," in *Proc. ICSE-SEIP*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.12944>
 - **[ParallelScience, Parallel Science]** Claude and the Denario Core Team, "The Parallel Science Project: Cyber space for human–AI co-evolution of science," *Parallel ArXiv*, PX: 2604.00017, Apr. 2026; the project also runs Parallel ArXiv (<https://papers.parallelspace.org>) and Parallel OpenReview (<https://reviews.parallelspace.org>), with code at <https://github.com/ParallelScience> and a landing page at <https://parallelspace.org>. [Online]. Available: <https://papers.parallelspace.org/abs/2604.00017>
 - **[Park et al., Generative agents]** J. S. Park, J. C. O'Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein, "Generative agents: Interactive simulacra of human behavior," in *Proc. UIST*, 2023. [Online]. Available: <https://arxiv.org/abs/2304.03442>
 - **[Park et al., OffsetBias]** J. Park, *et al.*, "OffsetBias: Leveraging debiased data for tuning evaluators," in *Findings of EMNLP*, 2024. [Online]. Available: <https://arxiv.org/abs/2407.06551>
 - **[Pathak et al., Detecting silent failures]** D. Pathak, H. Kumar, *et al.*, "Detecting silent failures in multi-agentic AI trajectories (work in progress paper)," in *Proc. ICPE Companion*, pp. 33–39, 2026. [Online]. Available: <https://arxiv.org/abs/2511.04032>
 - **[Patil et al., BFCL]** S. G. Patil, H. Mao, F. Yan, *et al.*, "The Berkeley Function Calling Leaderboard (BFCL): From tool use to agentic evaluation of large language models," in *Proc. ICML*, 2025, pp. 48371–48392. [Online]. Available: <https://proceedings.mlr.press/v267/patil25a.html>
 - **[Patil et al., Gorilla]** S. G. Patil, T. Zhang, X. Wang, and J. E. Gonzalez, "Gorilla: Large language model connected with massive APIs," in *Proc. NeurIPS*, 2024. [Online]. Available: <https://arxiv.org/abs/2305.15334>
 - **[Patwardhan et al., GDPval]** T. Patwardhan, *et al.* (OpenAI), "GDPval: Evaluating AI model performance on real-world economically valuable tasks," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2510.04374>

- **[Penfield, LoCoMo Audit]** Penfield Labs, "We audited LoCoMo: 6.4% of the answer key is wrong," *Penfield Labs*, Apr. 2026. [Online]. Available: <https://penfieldlabs.substack.com/p/we-audited-locomo-64-of-the-answer>
- **[Peng et al., Morae]** Y.-H. Peng, D. Li, J. P. Bigham, and A. Pavel, "Morae: Proactively pausing UI agents for user choices," in *Proc. ACM UIST*, 2025. [Online]. Available: <https://doi.org/10.1145/3746059.3747797>
- **[Peralta et al., Merged or Rejected]** S. R. O. Peralta, F. Hoshi, H. Washizaki, *et al.*, "Why are agentic pull requests merged or rejected? An empirical study," in *Proc. 23rd Int. Conf. Mining Software Repositories (MSR)*, Apr. 2026. [Online]. Available: <https://arxiv.org/abs/2605.22534>
- **[Plaat et al., Agentic large language]** A. Plaat, *et al.*, "Agentic large language models, a survey," *J. Artif. Intell. Res.*, vol. 84, Dec. 2025. [Online]. Available: <https://arxiv.org/abs/2503.23037>
- **[Portkey, Portkey AI Gateway]** Portkey, "Portkey AI Gateway — routing, caching, rate limiting, and budgets for LLM applications," *Portkey Documentation*, continuously updated. [Online]. Available: <https://portkey.ai/docs>
- **[Postman, State of the API 2025]** Postman, "2025 State of the API Report," survey of more than 5,700 developers, architects, and executives, 2025. [Online]. Available: <https://www.postman.com/state-of-api/2025/>
- **[Prasad et al., ADaPT]** A. Prasad, A. Koller, M. Hartmann, P. Clark, A. Sabharwal, M. Bansal, and T. Khot, "ADaPT: As-needed decomposition and planning with language models," in *Proc. NAACL Findings*, 2024. [Online]. Available: <https://arxiv.org/abs/2311.05772>
- **[Public agent incident canon]** Aggregated canon of public agent incident reports: Replit production-database deletion (Jul. 2025); Cursor `rm -rf` (Dec. 2025); Amazon Kiro Cost Explorer outage (Dec. 2025); Claude Cowork family-photo deletion (Feb. 2026). See per-incident references in this appendix.
- **[PyPI, LiteLLM Incident]** Python Package Index, "Incident Report: LiteLLM/Telnyx supply-chain attacks, with guidance," *The Python Package Index Blog*, Apr. 2026. [Online]. Available: <https://blog.pypi.org/posts/2026-04-02-incident-report-litellm-telnyx-supply-chain-attack/>
- **[Qi et al., CivRealm]** S. Qi, *et al.*, "CivRealm: A learning and reasoning odyssey in Civilization for decision-making agents," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2401.10568>

- **[Qi et al., MAD Amplified]** S. Qi, Y. Zou, P. Li, Z. Lin, X. Cheng, and D. Yu, "Amplified vulnerabilities: Structured jailbreak attacks on LLM-based multi-agent debate," *arXiv:2504.16489*, Apr. 2025. [Online]. Available: <https://arxiv.org/abs/2504.16489>
- **[Qian et al., ChatDev]** C. Qian, *et al.*, "ChatDev: Communicative agents for software development," in *Proc. ACL*, 2023 (ACL 2024). [Online]. Available: <https://arxiv.org/abs/2307.07924>
- **[Qin et al., ToolLLM]** Y. Qin, S. Liang, Y. Ye, *et al.*, "ToolLLM: Facilitating large language models to master 16000+ real-world APIs," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2307.16789>
- **[Rabanser et al., Agent Reliability]** S. Rabanser, *et al.*, "Towards a science of AI agent reliability," in *Proc. ICML*, 2026. [Online]. Available: <https://arxiv.org/abs/2602.16666>
- **[RabbitMQ Docs, RabbitMQ]** Broadcom / RabbitMQ Authors, "RabbitMQ — open-source message broker supporting multiple protocols, queues, and exchanges," *rabbitmq.com*, continuously updated. [Online]. Available: <https://www.rabbitmq.com/docs>
- **[Rafailov et al., DPO]** R. Rafailov, A. Sharma, E. Mitchell, S. Ermon, C. D. Manning, and C. Finn, "Direct preference optimization: Your language model is secretly a reward model," in *Proc. NeurIPS*, 2023. [Online]. Available: <https://arxiv.org/abs/2305.18290>
- **[Rajasekaran, Harness design]** P. Rajasekaran (Anthropic), "Harness design for long-running application development," *Anthropic Engineering*, Mar. 2026. [Online]. Available: <https://www.anthropic.com/engineering/harness-design-long-running-apps>
- **[Rank et al., PostTrainBench]** B. Rank, H. Bhatnagar, A. Prabhu, *et al.*, "PostTrainBench: Can LLM agents automate LLM post-training?," in *Proc. ICML*, 2026. [Online]. Available: <https://arxiv.org/abs/2603.08640>
- **[Rasmussen et al., Zep]** P. Rasmussen, *et al.*, "Zep: A temporal knowledge graph architecture for agent memory," *arXiv:2501.13956*, Jan. 2025. [Online]. Available: <https://arxiv.org/abs/2501.13956>
- **[Rawles et al., AndroidWorld]** C. Rawles, S. Clinckemaillie, Y. Chang, *et al.*, "AndroidWorld: A dynamic benchmarking environment for autonomous agents," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2405.14573>

- **[Raza et al., TRiSM for agentic]** S. Raza, *et al.*, "TRiSM for agentic AI: A review of trust, risk, and security management in LLM-based agentic multi-agent systems," *AI Open*, 2025. [Online]. Available: <https://arxiv.org/abs/2506.04133>
- **[Reason, Swiss Cheese]** J. Reason, "Human error: Models and management (Swiss cheese model)," *BMJ*, vol. 320, no. 7237, pp. 768–770, Mar. 2000. [Online]. Available: <https://pmc.ncbi.nlm.nih.gov/articles/PMC1117770/>
- **[Red Hat, KEDA vLLM]** Red Hat, "How to set up KServe autoscaling for vLLM with KEDA," *Red Hat Developer*, Sep. 2025. [Online]. Available: <https://developers.redhat.com/articles/2025/09/23/how-set-kserve-autoscaling-vllm-keda>
- **[Redis Docs, Redis documentation]** Redis Ltd., "Redis documentation — sorted sets, Lua scripting, and key expiration," *redis.io*, continuously updated. [Online]. Available: <https://redis.io/docs/>
- **[Redis Streams, Redis Streams]** Redis Ltd., "Redis Streams — append-only log with consumer groups, per-message acknowledgement, and replay," *redis.io*, continuously updated. [Online]. Available: <https://redis.io/docs/latest/develop/data-types/streams/>
- **[Rehberger, Copilot RCE]** J. Rehberger, "GitHub Copilot: Remote code execution via prompt injection (CVE-2025-53773)," *Embrace The Red*, 12 Aug. 2025. The researcher's own write-up: injected instructions make Copilot write "chat.tools.autoApprove": true into `.vscode/settings.json`, disabling user confirmations. [Online]. Available: <https://embracethered.com/blog/posts/2025/github-copilot-remote-code-execution-via-prompt-injection/>
- **[Replit incident 2025]** AI Incident Database, "Incident 1152: LLM-Driven Replit Agent Reportedly Executed Unauthorized Destructive Commands During Code Freeze, Leading to Loss of Production Data," Jul. 2025; contemporaneous reporting in *The Register*, 21 Jul. 2025, and the CEO's response in *The Register*, 22 Jul. 2025, including "rolling out automatic DB dev/prod separation." Replit published no postmortem; its CEO responded via press and X. [Online]. Available: <https://incidentdatabase.ai/cite/1152/>
- **[Restate Docs, Restate]** Restate, "Restate — durable execution engine for distributed applications," *restate.dev*, continuously updated. [Online]. Available: <https://docs.restate.dev/>

- **[ReversingLabs, nullifAI]** K. Zanki (ReversingLabs), "Malicious ML models discovered on Hugging Face platform," *ReversingLabs Blog*, Feb. 2025. [Online]. Available: <https://www.reversinglabs.com/blog/r1-identifies-malware-ml-model-hosted-on-hugging-face>
- **[Richens et al., General Agents]** J. Richens, T. Everitt, and D. Abel, "General agents need world models," in *Proc. ICML*, 2025. [Online]. Available: <https://proceedings.mlr.press/v267/richens25a.html>
- **[Riemer et al., Collusion Certification]** M. Riemer, T. Tosato, A. Memarian, *et al.*, "Position: Collusion risks among AI reasoning agents justify certification requirements for making market decisions," in *Proc. ICML (Position Paper Track)*, 2026. [Online]. Available: <https://arxiv.org/abs/2608.18078>
- **[Roberts et al., Data contamination through]** M. Roberts, *et al.*, "To the cutoff... and beyond? A longitudinal perspective on LLM data contamination," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2310.10628>
- **[Rootly AI Labs, SRE-skills-bench]** Rootly AI Labs, "SRE-skills-bench: LLM benchmark for SRE tasks," *sreskillsbench.com*, 2025. [Online]. Available: <https://sreskillsbench.com/>
- **[Rosenblatt, Perceptron]** F. Rosenblatt, "The Perceptron: A probabilistic model for information storage and organization in the brain," *Psychological Review*, vol. 65, no. 6, pp. 386–408, 1958. [Online]. Available: <https://doi.org/10.1037/h0042519>
- **[Rosser & Foerster, AgentBreeder]** J. Rosser and J. Foerster, "AgentBreeder: Mitigating the AI safety risks of multi-agent scaffolds via self-improvement," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2502.00757>
- **[Rumelhart et al., Backpropagation]** D. E. Rumelhart, G. E. Hinton, and R. J. Williams, "Learning representations by back-propagating errors," *Nature*, vol. 323, no. 6088, pp. 533–536, Oct. 1986. [Online]. Available: <https://doi.org/10.1038/323533a0>
- **[Russell & Norvig, AIMA]** S. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*, 4th ed. Hoboken, NJ: Pearson, 2021. [Online]. Available: <https://www.pearson.com/en-us/subject-catalog/p/artificial-intelligence-a-modern-approach/P200000003500>
- **[Sainz et al., NLP evaluation]** O. Sainz, *et al.*, "NLP evaluation in trouble: On the need to measure LLM data contamination for each benchmark," in *Proc. EMNLP Findings*, 2023. [Online]. Available: <https://arxiv.org/abs/2310.18018>

- **[Sakana AI, AI Scientist]** Sakana AI, "The AI Scientist: Towards fully automated open-ended scientific discovery," *Sakana AI Blog*, Aug. 2024. [Online]. Available: <https://sakana.ai/ai-scientist/>
- **[Sakshi et al., MMAU]** S. Sakshi, *et al.*, "MMAU: A massive multi-task audio understanding and reasoning benchmark," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.19168>
- **[Salesforce, Agentforce Pricing]** Salesforce, "Salesforce Agentforce pricing," *Salesforce*, continuously updated; checked Sep. 2026. [Online]. Available: <https://www.salesforce.com/agentforce/pricing/>
- **[Salim et al., Tokenomics]** M. Salim, *et al.*, "Tokenomics: Quantifying where tokens are used in agentic software engineering," *arXiv:2601.14470*, Jan. 2026. [Online]. Available: <https://arxiv.org/abs/2601.14470>
- **[Scale AI, RL Environments]** Scale AI, "Reinforcement learning environments for AI agents," *Scale AI*, continuously updated; checked Sep. 2026. [Online]. Available: <https://scale.com/rlenvironments>
- **[Scale AI, SWE-Bench Pro V2]** S. Dan, J. Da, *et al.* (Scale AI), "SWE-Bench Pro V2: A cleaner, harder-to-game leaderboard," *Scale Labs*, Sep. 2026. [Online]. Available: <https://labs.scale.com/blog/swe-bench-pro-v2>
- **[Schick et al., Toolformer]** T. Schick, J. Dwivedi-Yu, R. Dessì, *et al.*, "Toolformer: Language models can teach themselves to use tools," in *Proc. NeurIPS*, 2023. [Online]. Available: <https://arxiv.org/abs/2302.04761>
- **[SEC, Rule 204-2]** U.S. Securities and Exchange Commission, "Books and records to be maintained by investment advisers, 17 CFR 275.204-2(a)(17)(i)," *Code of Federal Regulations*, Title 17. [Online]. Available: <https://www.ecfr.gov/current/title-17/chapter-II/part-275/section-275.204-2>
- **[Seo et al., SimuHome]** G. Seo, J. Yang, J. Pyo, N. Kim, J. Lee, and Y. Jo, "SimuHome: A temporal- and environment-aware benchmark for smart home LLM agents," in *Proc. ICLR*, 2026. [Online]. Available: <https://openreview.net/forum?id=LCS1WsGvha>
- **[SEP-1865]** I. Salomon, L. Yosef, O. Chafik, *et al.*, "SEP-1865: MCP Apps: Interactive user interfaces for MCP," *Model Context Protocol extension specification*, stable Jan. 2026. [Online]. Available: <https://github.com/modelcontextprotocol/ext-apps/blob/main/specification/2026-01-26/apps.mdx>

- **[SEP-414]** ModelContextProtocol Authors, "SEP-414: MCP trace-context propagation via params._meta," 2026. [Online]. Available: <https://github.com/modelcontextprotocol/modelcontextprotocol/pull/414>
- **[SGLang, RadixAttention]** SGLang Team, "SGLang — a fast serving framework for large language and vision-language models (RadixAttention)," *docs.sglang.io*, continuously updated. [Online]. Available: <https://docs.sglang.io/>
- **[Shao et al., GRPO]** Z. Shao, P. Wang, Q. Zhu, *et al.*, "DeepSeekMath: Pushing the limits of mathematical reasoning in open language models (GRPO)," *arXiv:2402.03300*, Feb. 2024. [Online]. Available: <https://arxiv.org/abs/2402.03300>
- **[Shao et al., Misevolution]** S. Shao, Q. Ren, D. Liu, *et al.*, "Your agent may misevolve: Emergent risks in self-evolving LLM agents," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2509.26354>
- **[Shao et al., PrivacyLens]** Y. Shao, T. Li, W. Shi, Y. Liu, and D. Yang, "PrivacyLens: Evaluating privacy norm awareness of language models in action," in *Proc. NeurIPS Datasets and Benchmarks Track*, 2024. [Online]. Available: <https://arxiv.org/abs/2409.00138>
- **[Sharma et al., Constitutional classifiers]** M. Sharma, *et al.*, "Constitutional classifiers: Defending against universal jailbreaks across thousands of hours of red teaming," *arXiv:2501.18837*, Jan. 2025. [Online]. Available: <https://arxiv.org/abs/2501.18837>
- **[Shazeer et al., MoE]** N. Shazeer, A. Mirhoseini, K. Maziarz, *et al.*, "Outrageously large neural networks: The sparsely-gated mixture-of-experts layer," in *Proc. ICLR*, 2017. [Online]. Available: <https://arxiv.org/abs/1701.06538>
- **[Shen et al., Test-Time Interaction]** J. Shen, H. Bai, L. Zhang, *et al.*, "Thinking vs. doing: Improving agent reasoning by scaling test-time interaction," in *Proc. NeurIPS*, 2025. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2025/hash/f7c4783621a20f5f11a316cd249e0252-Abstract-Conference.html
- **[Sheng et al., S-LoRA]** Y. Sheng, S. Cao, D. Li, C. Hooper, N. Lee, S. Yang, C. Chou, B. Zhu, L. Zheng, K. Keutzer, J. E. Gonzalez, and I. Stoica, "S-LoRA: Serving thousands of concurrent LoRA adapters," *arXiv:2311.03285*, Nov. 2023. [Online]. Available: <https://arxiv.org/abs/2311.03285>
- **[Shi et al., Knowledge Transfer]** Q. Shi, C. E. Jimenez, S. Yao, *et al.*, "When models know more than they can explain: Quantifying knowledge transfer in human-AI collaboration," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://>

proceedings.neurips.cc/paper_files/paper/2025/hash/
975d11c51406cd10be48e47b36fb8698-Abstract-Conference.html

- **[Shinn et al., Reflexion]** N. Shinn, F. Cassano, E. Berman, A. Gopinath, K. Narasimhan, and S. Yao, "Reflexion: Language agents with verbal reinforcement learning," in *Proc. NeurIPS*, 2023. [Online]. Available: <https://arxiv.org/abs/2303.11366>
- **[Shortliffe, MYCIN]** E. H. Shortliffe, *Computer-Based Medical Consultations: MYCIN*. New York: Elsevier/North-Holland, 1976, ISBN 9780444001795. [Online]. Available: <https://archive.org/details/computerbasedmed0000shor>
- **[Silver et al., AlphaGo]** D. Silver, A. Huang, C. J. Maddison, *et al.*, "Mastering the game of Go with deep neural networks and tree search (AlphaGo)," *Nature*, vol. 529, no. 7587, pp. 484–489, Jan. 2016. [Online]. Available: <https://doi.org/10.1038/nature16961>
- **[Silver et al., AlphaZero]** D. Silver, T. Hubert, J. Schrittwieser, *et al.*, "A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play (AlphaZero)," *Science*, vol. 362, no. 6419, pp. 1140–1144, Dec. 2018. [Online]. Available: <https://doi.org/10.1126/science.aar6404>
- **[Simons Foundation, Meet Denario]** Simons Foundation, "Meet Denario, an AI assistant for every step of the scientific process," *Simons Foundation News*, 4 Nov. 2025; the Flatiron Institute team's own account of testing Denario end to end hundreds of times across twelve disciplines, with about 10 percent of outputs raising an intriguing question or finding. [Online]. Available: <https://www.simonsfoundation.org/2025/11/04/meet-denario-an-ai-assistant-for-every-step-of-the-scientific-process/>
- **[Sinha et al., Long-Horizon Execution]** A. Sinha, A. Arun, S. Goel, S. Staab, and J. Geiping, "The illusion of diminishing returns: Measuring long horizon execution in LLMs," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2509.09677>
- **[Skeptical, Skeptical]** Empyrean Labs, "Skeptical: the AI reviewer and verifier," *skeptical-ai.org*, 2026. [Online]. Available: <https://skeptical-ai.org>
- **[Smith, Contract Net Protocol]** R. G. Smith, "The Contract Net Protocol: High-level communication and control in a distributed problem solver," *IEEE Transactions on Computers*, vol. C-29, no. 12, pp. 1104–1113, Dec. 1980. [Online]. Available: <https://doi.org/10.1109/TC.1980.1675516>
- **[Smithery incident 2025]** GitGuardian, "From path traversal to supply chain compromise: Breaking MCP server hosting," *GitGuardian Blog*, reported and patched

Jun. 2025; publicly disclosed Oct. 2025. [Online]. Available: <https://blog.gitguardian.com/breaking-mcp-server-hosting/>

- **[Snorkel AI, Senior SWE-bench]** H. R. Ehrenberg, V. S. Chen, A. W. Hanjie, G. Orlanski, F. Sala, and K. R. Narasimhan (Snorkel AI, Princeton University, and University of Wisconsin-Madison), "Senior SWE-bench: Evaluating agents on their ability to act as senior engineers," *Snorkel AI*, Jul. 2026. [Online]. Available: <https://senior-swe-bench.snorkel.ai/>
- **[Solomon et al., LumiMAS]** R. Solomon, Y. Yerushalmi Levi, *et al.*, "LumiMAS: A comprehensive framework for real-time monitoring and enhanced observability in multi-agent systems," in *Proc. AAMAS*, 2025. [Online]. Available: <https://arxiv.org/abs/2508.12412>
- **[South et al., Authenticated Delegation]** T. South, S. Marro, T. Hardjono, *et al.*, "Position: AI agents need authenticated delegation," in *Proc. ICML*, 2025. [Online]. Available: <https://proceedings.mlr.press/v267/south25a.html>
- **[Splunk, Agent Control]** Splunk (Cisco), "Agent Control: centralized evaluation of LLM and tool inputs and outputs during agent execution," *Splunk Agent Observability documentation*, continuously updated. [Online]. Available: <https://agent-observability-docs.splunk.com/concepts/agent-control/overview>
- **[Splunk, AI Infrastructure Monitoring]** Splunk (Cisco), "Set up AI infrastructure monitoring: built-in navigators and dashboards for LLM services and gateways, model-serving platforms, vector databases, and the infrastructure beneath them, collected through the Splunk Distribution of the OpenTelemetry Collector," *Splunk Observability Cloud documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://help.splunk.com/en/splunk-observability-cloud/observability-for-ai/splunk-ai-infrastructure-monitoring>
- **[Splunk, Luna]** Splunk (Cisco), "Luna-2: small language models for AI evaluation," *Splunk Agent Observability documentation*, continuously updated. [Online]. Available: <https://agent-observability-docs.splunk.com/concepts/luna/luna>
- **[Splunk, Signals]** Splunk (Cisco), "Signals: proactive discovery of failure patterns across agent traces," *Splunk Agent Observability documentation*, continuously updated. [Online]. Available: <https://agent-observability-docs.splunk.com/concepts/signals>

- **[Sprague et al., To CoT or Not]** Z. Sprague, F. Yin, J. D. Rodriguez, *et al.*, "To CoT or not to CoT? Chain-of-thought helps mainly on math and symbolic reasoning," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2409.12183>
- **[Squire, Mechanisms of memory]** L. R. Squire, "Mechanisms of memory," *Science*, vol. 232, no. 4758, pp. 1612–1619, Jun. 1986. [Online]. Available: <https://doi.org/10.1126/science.3086978>
- **[Srivastava et al., BIG-bench]** A. Srivastava, *et al.*, "Beyond the imitation game (BIG-bench): Quantifying and extrapolating the capabilities of language models," *Trans. Mach. Learn. Res.*, 2022. [Online]. Available: <https://arxiv.org/abs/2206.04615>
- **[Strathern, Improving Ratings]** M. Strathern, "'Improving ratings': Audit in the British university system," *European Review*, vol. 5, no. 3, pp. 305–321, Jul. 1997. [Online]. Available: [https://doi.org/10.1002/\(SICI\)1234-981X\(199707\)5:3%3C305::AID-EUR0184%3E3.0.CO;2-4](https://doi.org/10.1002/(SICI)1234-981X(199707)5:3%3C305::AID-EUR0184%3E3.0.CO;2-4)
- **[Stripe Idempotency]** Stripe, "API: idempotent requests," *Stripe API Reference*, continuously updated. [Online]. Available: https://docs.stripe.com/api/idempotent_requests
- **[Su et al., LLM-Friendly Knowledge]** H. Su, W. Luo, W. Han, *et al.*, "LLM-friendly knowledge representation for customer support," in *Proc. 31st Int. Conf. Computational Linguistics: Industry Track*, 2025, pp. 496–504. [Online]. Available: <https://aclanthology.org/2025.coling-industry.42/>
- **[Subramaniam et al., Multiagent Finetuning]** V. Subramaniam, Y. Du, J. B. Tenenbaum, *et al.*, "Multiagent finetuning: Self improvement with diverse reasoning chains," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2501.05707>
- **[Sun et al., Agents' Last Exam]** Y. Sun, X. Han, W. Zhang, *et al.*, "Agents' Last Exam," in *Proc. NeurIPS Evaluations and Datasets Track*, 2026. [Online]. Available: <https://arxiv.org/abs/2606.05405>
- **[Sun et al., PathGen-1.6M]** Y. Sun, Y. Zhang, Y. Si, *et al.*, "PathGen-1.6M: 1.6 million pathology image-text pairs generation through multi-agent collaboration," in *Proc. ICLR*, 2025. [Online]. Available: <https://openreview.net/forum?id=rFpZnn11gj>
- **[Sutton, Bitter Lesson]** R. S. Sutton, "The bitter lesson," *incompleteideas.net*, Mar. 2019. [Online]. Available: <http://www.incompleteideas.net/IncIdeas/BitterLesson.html>

- **[SWE-Bench Pro]** X. Deng, *et al.* (Scale AI), "SWE-Bench Pro: Can AI agents solve long-horizon software engineering tasks?" in *Proc. ICML*, 2026. [Online]. Available: <https://arxiv.org/abs/2509.16941>
- **[Sánchez-Torrón et al., Write or Automate]** M. Sánchez-Torrón, D. Akselrod, and J. Rauchwerk, "To write or to automate linguistic prompts, that is the question," in *Proc. 26th Annu. Conf. European Assoc. Machine Translation (Volume 1)*, pp. 85–95, Jun. 2026. [Online]. Available: <https://aclanthology.org/2026.eamt-1.8/>
- **[Tam et al., Speak Freely]** Z. R. Tam, *et al.*, "Let me speak freely? A study on the impact of format restrictions on performance of large language models," in *Proc. EMNLP Industry Track*, pp. 1218–1236, Nov. 2024. [Online]. Available: <https://arxiv.org/abs/2408.02442>
- **[Tan et al., JudgeBench]** S. Tan, S. Zhuang, K. Montgomery, *et al.*, "JudgeBench: A benchmark for evaluating LLM-based judges," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.12784>
- **[Tang et al., Coding Agents Fail]** N. Tang, C. Chen, G. Xu, *et al.*, "How coding agents fail their users: A large-scale analysis of developer-agent misalignment in 20,574 real-world sessions," in *Proc. EMNLP*, 2026. [Online]. Available: <https://arxiv.org/abs/2605.29442>
- **[Tang et al., Dark Patterns and GUI Agents]** J. Tang, C. Chen, J. Li, *et al.*, "Dark patterns meet GUI agents: LLM agent susceptibility to manipulative interfaces and the role of human oversight," in *Proc. ACM CHI*, 2026. [Online]. Available: <https://doi.org/10.1145/3772318.3791568>
- **[TechCrunch, Taco Bell AI]** A. Ha, "Taco Bell is having second thoughts about relying on AI at the drive-through," *TechCrunch*, Aug. 2025. [Online]. Available: <https://techcrunch.com/2025/08/30/taco-bell-is-having-second-thoughts-about-relying-on-ai-at-the-drive-through/>
- **[Temporal Docs, Temporal]** Temporal Technologies, "Temporal — durable execution platform for workflows as code," *docs.temporal.io*, continuously updated. [Online]. Available: <https://docs.temporal.io/>
- **[Temporal, Durable AI]** Temporal Technologies, "Durable execution meets AI: Why Temporal is the perfect foundation for AI agents," *Temporal Blog*, 2025. [Online]. Available: <https://temporal.io/blog/durable-execution-meets-ai-why-temporal-is-the-perfect-foundation-for-ai>

- **[Temporal, Dynamic Agents]** M. Egger and S. Androulakis (Temporal), "Of course you can build dynamic AI agents with Temporal," *Temporal Blog*, Nov. 2025. [Online]. Available: <https://temporal.io/blog/of-course-you-can-build-dynamic-ai-agents-with-temporal>
- **[Terminal-Bench, Terminal-Bench]** Terminal-Bench Team, "Terminal-Bench 2.1: A revision of Terminal-Bench 2.0 that fixes 28 tasks," *Terminal-Bench*, May 2026. [Online]. Available: <https://www.tbench.ai/news/terminal-bench-2-1>
- **[Terminal-Bench, Terminal-Bench 4.0]** Terminal-Bench Team, "Terminal-Bench 4.0: Calibrating task resources, fixing tasks, and removing saturated tasks," *Terminal-Bench*, Aug. 2026. [Online]. Available: <https://www.tbench.ai/news/terminal-bench-4-0>
- **[Terminal-Bench-Science, Terminal-Bench-Science 0.1]** Terminal-Bench-Science Team (Stanford University and Laude Institute), "Terminal-Bench-Science 0.1: Evaluating agents on real computational workflows from scientific research," *terminal-bench-science.ai*, Aug. 2026. [Online]. Available: <https://www.terminal-bench-science.ai/announcement>
- **[Tomasev et al., Virtual Agent Economies]** N. Tomasev, M. Franklin, J. Z. Leibo, J. Jacobs, W. A. Cunningham, I. Gabriel, and S. Osindero (Google DeepMind), "Virtual agent economies," *arXiv:2509.10147*, Sep. 2025. [Online]. Available: <https://arxiv.org/abs/2509.10147>
- **[Trace-store incident canon]** Aggregated canon of public trace-store breaches: DeepSeek ClickHouse exposure (Wiz Research, Jan. 2025); ChatGPT URL leak (Nov. 2025); OpenAI / Mixpanel breach (Nov. 2025). [Online]. Available: <https://www.wiz.io/blog/wiz-research-uncovers-exposed-deepseek-database-leak>
- **[Trail of Bits, Context Protector]** Trail of Bits, "We built the security layer MCP always needed (mcp-context-protector, trust-on-first-use pinning)," *Trail of Bits Blog*, Jul. 2025. [Online]. Available: <https://blog.trailofbits.com/2025/07/28/we-built-the-security-layer-mcp-always-needed/>
- **[Trail of Bits, Line Jumping]** Trail of Bits, "Jumping the line: how MCP servers can attack you before you ever use them," *Trail of Bits Blog*, Apr. 2025. [Online]. Available: <https://blog.trailofbits.com/2025/04/21/jumping-the-line-how-mcp-servers-can-attack-you-before-you-ever-use-them/>
- **[Tran & Kiela, Single-agent LLMs outperform]** D. Tran and D. Kiela, "Single-agent LLMs outperform multi-agent systems on multi-hop reasoning under equal thinking

token budgets," *arXiv:2604.02460*, Apr. 2026. [Online]. Available: <https://arxiv.org/abs/2604.02460>

- **[Trilogy, Technical deep dive]** Trilogy AI, "Technical deep dive: Hermes vs. OpenClaw, two approaches to personal AI infrastructure," *Trilogy AI Substack*, 2026. [Online]. Available: <https://trilogyai.substack.com/p/technical-deep-dive-hermes-vs-openclaw>
- **[Tsai et al., RTLFixer]** Y.-D. Tsai, M. Liu, and H. Ren, "RTLFixer: Automatically fixing RTL syntax errors with large language models," in *Proc. 61st ACM/IEEE Design Automation Conf. (DAC)*, Jun. 2024. [Online]. Available: <https://arxiv.org/abs/2311.16543>
- **[Tulving, Episodic and semantic]** E. Tulving, "Episodic and semantic memory," in *Organization of Memory*, E. Tulving and W. Donaldson, Eds. New York: Academic Press, 1972, pp. 381–403.
- **[UK AISI, Inspect]** UK AI Security Institute and Meridian Labs, "Inspect: an open-source framework for large language model evaluations," *Inspect documentation*, continuously updated. [Online]. Available: <https://inspect.aisi.org.uk/>
- **[Vaswani et al., Transformer]** A. Vaswani, N. Shazeer, N. Parmar, *et al.*, "Attention is all you need," in *Proc. NeurIPS*, 2017. [Online]. Available: <https://arxiv.org/abs/1706.03762>
- **[Vercel, Removed Tools]** A. Qu (Vercel), "We removed 80% of our agent's tools," *Vercel Blog*, Dec. 2025. [Online]. Available: <https://vercel.com/blog/we-removed-80-percent-of-our-agents-tools>
- **[Vero et al., BaxBench]** M. Vero, N. Mündler, V. Chibotaru, *et al.*, "BaxBench: Can LLMs generate correct and secure backends?" in *Proc. ICML*, 2025. [Online]. Available: <https://arxiv.org/abs/2502.11844>
- **[Vertex AI Agent Builder, Vertex AI Agent]** Google Cloud, "Gemini Enterprise Agent Platform (formerly Vertex AI, including Vertex AI Agent Builder) — managed platform for building and deploying enterprise-grade AI agents," *Google Cloud Documentation*, continuously updated; checked Sep. 2026. [Online]. Available: <https://cloud.google.com/products/gemini-enterprise-agent-platform>
- **[Villaescusa-Navarro et al., Denario project]** F. Villaescusa-Navarro, *et al.*, "The Denario project: Deep knowledge AI agents for scientific discovery," *arXiv:2510.26887*, Oct. 2025. [Online]. Available: <https://arxiv.org/abs/2510.26887>

- **[vLLM Docs, vLLM]** vLLM Project, "vLLM — high-throughput and memory-efficient LLM inference and serving engine," *docs.vllm.ai*, continuously updated. [Online]. Available: <https://docs.vllm.ai/>
- **[VPI-Bench, VPI-Bench]** T. Cao, B. Lim, Y. Liu, *et al.*, "VPI-Bench: Visual prompt injection attacks for computer-use agents," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2506.02456>
- **[W3C Trace Context]** W3C, "Trace Context specification," W3C Recommendation, Nov. 2021. [Online]. Available: <https://www.w3.org/TR/trace-context/>
- **[W3C, DID]** W3C, "Decentralized Identifiers (DIDs) v1.0," W3C Recommendation, Jul. 2022. [Online]. Available: <https://www.w3.org/TR/did-core/>
- **[W3C, VC]** W3C, "Verifiable Credentials Data Model v2.0," W3C Recommendation, May 2025. [Online]. Available: <https://www.w3.org/TR/vc-data-model/>
- **[Wang & Chen, MIRIX]** Y. Wang and X. Chen, "MIRIX: Multi-agent memory system for LLM-based agents," *arXiv:2507.07957*, Jul. 2025. [Online]. Available: <https://arxiv.org/abs/2507.07957>
- **[Wang et al., Agent workflow memory]** Z. Wang, J. Mao, D. Fried, and G. Neubig, "Agent workflow memory," in *Proc. ICML*, 2024. [Online]. Available: <https://arxiv.org/abs/2409.07429>
- **[Wang et al., AgentTaxo]** Q. Wang, *et al.*, "AgentTaxo: Dissecting and benchmarking token distribution of LLM multi-agent systems," *ICLR 2025 Workshop on Foundation Models in the Wild*, Mar. 2025. [Online]. Available: <https://openreview.net/forum?id=0ilbiYYIpC>
- **[Wang et al., Berkeley RDI]** H. Wang, Q. Mang, A. Cheung, K. Sen, and D. Song, "We scored 100% on AI benchmarks without solving a single problem," *Berkeley RDI Blog*, Apr. 2026. [Online]. Available: <https://rdi.berkeley.edu/blog/trustworthy-benchmarks/>
- **[Wang et al., Developer Practices]** Y. Wang, X. Xu, J. Chen, T. Bi, W. Gu, and Z. Zheng, "An empirical study of agent developer practices in AI agent frameworks," *arXiv:2512.01939*, Dec. 2025. [Online]. Available: <https://arxiv.org/abs/2512.01939>
- **[Wang et al., Fair Evaluators]** P. Wang, L. Li, L. Chen, *et al.*, "Large language models are not fair evaluators," in *Proc. ACL*, 2023. [Online]. Available: <https://arxiv.org/abs/2305.17926>
- **[Wang et al., Huxley-Gödel Machine]** W. Wang, P. Piekos, *et al.*, "Huxley-Gödel Machine: Human-level coding agent development by an approximation of the optimal

- self-improving machine," in *Proc. ICLR Oral*, 2026. [Online]. Available: <https://openreview.net/forum?id=T0EiEuh00L>
- **[Wang et al., Math-Shepherd]** P. Wang, L. Li, Z. Shao, *et al.*, "Math-Shepherd: Verify and reinforce LLMs step-by-step without human annotations," in *Proc. ACL*, 2024. [Online]. Available: <https://arxiv.org/abs/2312.08935>
 - **[Wang et al., MCPTox]** Z. Wang, Y. Gao, Y. Wang, *et al.*, "MCPTox: A benchmark for tool poisoning on real-world MCP servers," in *Proc. AAAI*, pp. 35811–35819, 2026. [Online]. Available: <https://arxiv.org/abs/2508.14925>
 - **[Wang et al., PandaLM]** Y. Wang, *et al.*, "PandaLM: An automatic evaluation benchmark for LLM instruction tuning optimization," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2306.05087>
 - **[Wang et al., Plan-and-Solve]** L. Wang, W. Xu, Y. Lan, *et al.*, "Plan-and-solve prompting: Improving zero-shot chain-of-thought reasoning by large language models," in *Proc. ACL*, 2023. [Online]. Available: <https://arxiv.org/abs/2305.04091>
 - **[Wang et al., SuperGLUE]** A. Wang, Y. Pruksachatkun, N. Nangia, *et al.*, "SuperGLUE: A stickier benchmark for general-purpose language understanding systems," in *Proc. NeurIPS*, 2019. [Online]. Available: <https://arxiv.org/abs/1905.00537>
 - **[Wang et al., Voyager]** G. Wang, Y. Xie, Y. Jiang, *et al.*, "Voyager: An open-ended embodied agent with large language models," *Trans. Mach. Learn. Res.*, 2023. [Online]. Available: <https://arxiv.org/abs/2305.16291>
 - **[Wang, Sutando]** C. Wang, "Sutando: a personal multi-agent fleet across voice, Discord, Telegram, iMessage, and a menu-bar input," *GitHub*, 2025. [Online]. Available: <https://github.com/sonichi/sutando>
 - **[Wei et al., Chain-of-Thought]** J. Wei, X. Wang, D. Schuurmans, *et al.*, "Chain-of-thought prompting elicits reasoning in large language models," in *Proc. NeurIPS*, 2022. [Online]. Available: <https://arxiv.org/abs/2201.11903>
 - **[WHATWG, Server-sent events]** WHATWG, "Server-sent events," *HTML Living Standard*, continuously updated. [Online]. Available: <https://html.spec.whatwg.org/multipage/server-sent-events.html>
 - **[White et al., Autonomic Architecture]** S. R. White, J. E. Hanson, I. Whalley, D. M. Chess, and J. O. Kephart, "An architectural approach to autonomic computing," in *Proc. International Conference on Autonomic Computing (ICAC)*, 2004, pp. 2–9. [Online]. Available: <https://doi.org/10.1109/ICAC.2004.1301340>

- **[Willison, Lethal Trifecta]** S. Willison, "The lethal trifecta for AI agents: private data, untrusted content, and external communication," *simonwillison.net*, Jun. 2025. [Online]. Available: <https://simonwillison.net/2025/Jun/16/the-lethal-trifecta/>
- **[Willison, MCP]** S. Willison, "Model Context Protocol has prompt injection security problems (tool rug pulls and shadowing)," *simonwillison.net*, Apr. 2025. [Online]. Available: <https://simonwillison.net/2025/Apr/9/mcp-prompt-injection/>
- **[Winograd, SHRDLU]** T. Winograd, "Procedures as a representation for data in a computer program for understanding natural language," MIT Artificial Intelligence Laboratory, Technical Report 235, Feb. 1971. [Online]. Available: <https://dspace.mit.edu/handle/1721.1/7095>
- **[Wooldridge & Jennings, Intelligent Agents]** M. Wooldridge and N. R. Jennings, "Intelligent agents: Theory and practice," *The Knowledge Engineering Review*, vol. 10, no. 2, pp. 115–152, Jun. 1995. [Online]. Available: <https://doi.org/10.1017/S0269888900008122>
- **[Wu et al., AutoGen]** Q. Wu, G. Bansal, J. Zhang, *et al.*, "AutoGen: Enabling next-gen LLM applications via multi-agent conversation," in *Proc. COLM*, 2024. [Online]. Available: <https://arxiv.org/abs/2308.08155>
- **[Wu et al., LongMemEval]** D. Wu, *et al.*, "LongMemEval: Benchmarking chat assistants on long-term interactive memory," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.10813>
- **[Xi et al., AgentGym-RL]** Z. Xi, J. Huang, C. Liao, *et al.*, "AgentGym-RL: An open-source framework to train LLM agents for long-horizon decision making via multi-turn RL," in *Proc. ICLR Oral*, 2026. [Online]. Available: <https://openreview.net/forum?id=ZgCCDwcGwn>
- **[Xi et al., AgentPRM]** Z. Xi, *et al.*, "AgentPRM: Process reward models for LLM agents via step-wise promise and progress," in *Proc. WWW*, 2026. [Online]. Available: <https://arxiv.org/abs/2511.08325>
- **[Xie et al., OSWorld]** T. Xie, D. Zhang, J. Chen, *et al.*, "OSWorld: Benchmarking multimodal agents for open-ended tasks in real computer environments," in *Proc. NeurIPS Datasets and Benchmarks Track*, 2024. [Online]. Available: <https://arxiv.org/abs/2404.07972>

- **[Xie et al., TravelPlanner]** J. Xie, *et al.*, "TravelPlanner: A benchmark for real-world planning with language agents," in *Proc. ICML*, 2024. [Online]. Available: <https://arxiv.org/abs/2402.01622>
- **[Xie et al., WhodunitBench]** J. Xie, *et al.*, "WhodunitBench: Evaluating large multimodal agents via murder mystery games," in *Proc. NeurIPS Datasets and Benchmarks Track*, 2024. [Online]. Available: <https://openreview.net/forum?id=qmvtDIfbmS>
- **[Xiong et al., Experience-Following]** Z. Xiong, Y. Lin, W. Xie, *et al.*, "How memory management impacts LLM agents: An empirical study of experience-following behavior," in *Proc. 64th Annu. Meeting Assoc. Computational Linguistics (Volume 1: Long Papers)*, pp. 623–645, Jul. 2026. [Online]. Available: <https://aclanthology.org/2026.acl-long.27/>
- **[Xu et al., A-MEM]** W. Xu, Z. Liang, *et al.*, "A-MEM: Agentic memory for LLM agents," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2502.12110>
- **[Xu et al., MedAgentGym]** R. Xu, Y. Zhuang, *et al.*, "MedAgentGym: A scalable agentic training environment for code-centric reasoning in biomedical data science," in *Proc. ICLR Oral*, 2026. [Online]. Available: <https://openreview.net/forum?id=jHDZEUGS4r>
- **[Xu et al., Planning & Control]** L. Xu, *et al.*, "Open-source planning & control system with language agents for autonomous scientific discovery," in *Proc. ICML Workshop on Machine Learning for Astrophysics*, 2025. [Online]. Available: <https://arxiv.org/abs/2507.07257>
- **[Xu et al., TheAgentCompany]** F. F. Xu, Y. Song, B. Li, *et al.*, "TheAgentCompany: Benchmarking LLM agents on consequential real world tasks," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2412.14161>
- **[Xu, Competing with AI Scientists]** L. Xu, "Competing with AI scientists: A fully agent-driven approach to cosmological parameter inference," solution talk, *NeurIPS 2025 Competition Track, FAIR Universe*, Dec. 2025. [Online]. Available: <https://neurips.cc/virtual/2025/136457>
- **[Yang et al., Agent Scaling]** Y. Yang, *et al.*, "Understanding agent scaling in LLM-based multi-agent systems via diversity," in *Proc. ICLR Workshop on AI for Mechanism Design and Strategic Decision Making (non-archival)*, 2026. [Online]. Available: <https://arxiv.org/abs/2602.03794>
- **[Yang et al., EmbodiedBench]** R. Yang, H. Chen, J. Zhang, *et al.*, "EmbodiedBench: Comprehensive benchmarking multi-modal large language models for vision-driven

- embodied agents," in *Proc. ICML*, 2025. [Online]. Available: <https://openreview.net/forum?id=DgGF2LEBPS>
- **[Yang et al., MAS Concurrency]** X. Yang, L. Li, Z. Ji, T. J. Zhang, and W. Jiang, "Position: Multi-agent systems should prioritize concurrency control," in *Proc. ICML*, 2026. [Online]. Available: <https://arxiv.org/abs/2608.18092>
 - **[Yang et al., SWE-agent]** J. Yang, C. E. Jimenez, A. Wettig, K. Lieret, S. Yao, K. Narasimhan, and O. Press, "SWE-agent: Agent-computer interfaces enable automated software engineering," in *Proc. NeurIPS*, 2024. [Online]. Available: <https://arxiv.org/abs/2405.15793>
 - **[Yao et al., CacheBlend]** J. Yao, H. Li, Y. Liu, *et al.*, "CacheBlend: Fast large language model serving for RAG with cached knowledge fusion," in *Proc. EuroSys*, 2025. [Online]. Available: <https://arxiv.org/abs/2405.16444>
 - **[Yao et al., ReAct]** S. Yao, J. Zhao, D. Yu, *et al.*, "ReAct: Synergizing reasoning and acting in language models," in *Proc. ICLR*, 2023. [Online]. Available: <https://arxiv.org/abs/2210.03629>
 - **[Yao et al., Retroformer]** W. Yao, *et al.*, "Retroformer: Retrospective large language agents with policy gradient optimization," in *Proc. ICLR Spotlight*, 2024. [Online]. Available: <https://arxiv.org/abs/2308.02151>
 - **[Yao et al., Sycophancy]** B. Yao, C. Shang, W. Du, J. He, R. Lian, Y. Zhang, H. Su, S. Swamy, and Y. Qi, "Peacemaker or troublemaker: How sycophancy shapes multi-agent debate," *arXiv:2509.23055*, Sep. 2025. [Online]. Available: <https://arxiv.org/abs/2509.23055>
 - **[Yao et al., Tree of Thoughts]** S. Yao, D. Yu, J. Zhao, *et al.*, "Tree of thoughts: Deliberate problem solving with large language models," in *Proc. NeurIPS*, 2023. [Online]. Available: <https://arxiv.org/abs/2305.10601>
 - **[Yao et al., T-bench]** S. Yao, N. Shinn, P. Razavi, and K. Narasimhan, "T-bench: A benchmark for tool-agent-user interaction in real-world domains," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2406.12045>
 - **[Yehudai et al., Survey on evaluation]** A. Yehudai, *et al.*, "Survey on evaluation of LLM-based agents," in *ACL Findings*, 2025. [Online]. Available: <https://arxiv.org/abs/2503.16416>
 - **[Yoshioka et al., Meaningful Pull Requests]** H. Yoshioka, T. Monno, H. Tokumasu, *et al.*, "Let's make every pull request meaningful: An empirical analysis of developer and agentic pull requests," in *Proc. 23rd Int. Conf. Mining Software Repositories (MSR)*, Mining

Challenge Track, Apr. 2026. [Online]. Available: <https://doi.org/10.1145/3793302.3793567>

- **[Young, Effective Harnesses]** J. Young (Anthropic), "Effective harnesses for long-running agents," *Anthropic Engineering*, Nov. 2025. [Online]. Available: <https://www.anthropic.com/engineering/effective-harnesses-for-long-running-agents>
- **[Yu et al., DAPO]** Q. Yu, Z. Zhang, R. Zhu, *et al.*, "DAPO: An open-source LLM reinforcement learning system at scale," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2503.14476>
- **[Yu et al., MemAgent]** H. Yu, T. Chen, *et al.*, "MemAgent: Reshaping long-context LLM with multi-conv RL-based memory agent," in *Proc. ICLR Oral*, 2026. [Online]. Available: <https://openreview.net/forum?id=k5nIOvYGCL>
- **[Yu et al., Multi-agent memory]** Z. Yu, N. Yu, *et al.*, "Multi-agent memory from a computer architecture perspective: Visions and challenges ahead," in *Proc. Architecture 2.0 Workshop at ASPLOS (non-archival)*, 2026. [Online]. Available: <https://arxiv.org/abs/2603.10062>
- **[Yu et al., MYCIN Evaluation]** V. L. Yu, L. M. Fagan, *et al.*, "Antimicrobial selection by a computer: A blinded evaluation by infectious diseases experts," *JAMA*, vol. 242, no. 12, pp. 1279–1282, Sep. 1979. [Online]. Available: <https://pubmed.ncbi.nlm.nih.gov/480542/>
- **[Yu et al., Orca]** G.-I. Yu, J. S. Jeong, G.-W. Kim, S. Kim, and B.-G. Chun, "Orca: A distributed serving system for transformer-based generative models (continuous batching)," in *Proc. OSDI*, 2022. [Online]. Available: <https://www.usenix.org/conference/osdi22/presentation/yu>
- **[Yuan et al., Numerical Nondeterminism]** J. Yuan, H. Li, X. Ding, *et al.*, "Understanding and mitigating numerical sources of nondeterminism in LLM inference," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2506.09501>
- **[Yuan et al., OSWorld 2.0]** M. Yuan, Z. Zhou, X. Xiong, *et al.* (XLANG Lab), "OSWorld 2.0: Benchmarking computer use agents on long-horizon real-world tasks," *arXiv: 2606.29537*, Jul. 2026. [Online]. Available: <https://arxiv.org/abs/2606.29537>
- **[Yuan et al., Skill Composition]** L. Yuan, W. Chen, Y. Zhang, *et al.*, "From $f(x)$ and $g(x)$ to $f(g(x))$: LLMs learn new skills in RL by composing old ones," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2509.25123>

- **[Yue et al., Does reinforcement learning]** Y. Yue, *et al.*, "Does reinforcement learning really incentivize reasoning capacity in LLMs beyond the base model?" in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2504.13837>
- **[Yuksekgonul et al., TextGrad]** M. Yuksekgonul, F. Bianchi, J. Boen, *et al.*, "TextGrad: Automatic 'differentiation' via text," *arXiv:2406.07496*, Jun. 2024. [Online]. Available: <https://arxiv.org/abs/2406.07496>
- **[Zhang et al., ACE]** Q. Zhang, *et al.*, "Agentic context engineering: Evolving contexts for self-improving language models," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2510.04618>
- **[Zhang et al., AFlow]** J. Zhang, J. Xiang, *et al.*, "AFlow: Automating agentic workflow generation," in *Proc. ICLR Oral*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.10762>
- **[Zhang et al., AgentOptimizer]** S. Zhang, *et al.*, "Offline training of language model agents with functions as learnable weights (AgentOptimizer)," in *Proc. ICML*, 2024. [Online]. Available: <https://arxiv.org/abs/2402.11359>
- **[Zhang et al., AgentPrune]** G. Zhang, Y. Yue, Z. Li, S. Yun, G. Wan, K. Wang, D. Cheng, J. X. Yu, and T. Chen, "Cut the crap: An economical communication pipeline for LLM-based multi-agent systems," in *Proc. ICLR*, 2025. [Online]. Available: https://proceedings.iclr.cc/paper_files/paper/2025/hash/bbc461518c59a2a8d64e70e2c38c4a0e-Abstract-Conference.html
- **[Zhang et al., AgenTracer]** G. Zhang, J. Wang, J. Chen, W. Zhou, K. Wang, and S. Yan, "AgenTracer: Who is inducing failure in the LLM agentic systems?," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2509.03312>
- **[Zhang et al., Cybench]** A. K. Zhang, N. Perry, *et al.*, "Cybench: A framework for evaluating cybersecurity capabilities and risks of language models," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2408.08926>
- **[Zhang et al., Darwin Gödel Machine]** J. Zhang, S. Hu, C. Lu, R. Lange, and J. Clune, "Darwin Gödel Machine: Open-ended evolution of self-improving agents," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2505.22954>
- **[Zhang et al., G-Designer]** G. Zhang, Y. Yue, X. Sun, *et al.*, "G-Designer: Architecting multi-agent communication topologies via graph neural networks," in *Proc. ICML Spotlight*, 2025. [Online]. Available: <https://openreview.net/forum?id=LpE54NUnmO>

- **[Zhang et al., G-Memory]** G. Zhang, *et al.*, "G-Memory: Tracing hierarchical memory for multi-agent systems," in *Proc. NeurIPS Spotlight*, 2025. [Online]. Available: <https://arxiv.org/abs/2506.07398>
- **[Zhang et al., landscape of agentic]** G. Zhang, *et al.*, "The landscape of agentic reinforcement learning for LLMs: A survey," *Transactions on Machine Learning Research (TMLR)*, 2026. [Online]. Available: <https://arxiv.org/abs/2509.02547>
- **[Zhang et al., MaAS]** G. Zhang, L. Niu, J. Fang, K. Wang, *et al.*, "Multi-agent architecture search via agentic supernet," in *Proc. ICML Oral*, 2025. [Online]. Available: <https://openreview.net/forum?id=imcyV1zpXh>
- **[Zhang et al., Memory Survey]** Z. Zhang, *et al.*, "A survey on the memory mechanism of large language model based agents," *ACM Trans. Inf. Syst.*, 2024. [Online]. Available: <https://arxiv.org/abs/2404.13501>
- **[Zhang et al., Who&When]** S. Zhang, *et al.*, "Which agent causes task failures and when? On automated failure attribution of LLM multi-agent systems (Who&When)," in *Proc. ICML Spotlight*, 2025. [Online]. Available: <https://arxiv.org/abs/2505.00212>
- **[Zhao et al., Absolute Zero]** A. Zhao, Y. Wu, Y. Yue, *et al.*, "Absolute Zero: Reinforced self-play reasoning with zero data," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2505.03335>
- **[Zhao et al., Agent-in-the-Loop]** C. Zhao, T. Zhang, H. Su, *et al.*, "Agent-in-the-loop: A data flywheel for continuous improvement in LLM-based customer support," in *Proc. 2025 Conf. Empirical Methods in Natural Language Processing: Industry Track*, 2025, pp. 1919–1930. [Online]. Available: <https://aclanthology.org/2025.emnlp-industry.135/>
- **[Zhao et al., CompeteAI]** Q. Zhao, *et al.*, "CompeteAI: Understanding the competition dynamics of large language model-based agents," in *Proc. ICML*, 2024. [Online]. Available: <https://arxiv.org/abs/2310.17512>
- **[Zhao et al., Stronger-MAS]** Y. Zhao, L. Hu, Y. Wang, *et al.*, "Stronger-MAS: Multi-agent reinforcement learning for collaborative LLMs," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2510.11062>
- **[Zheng et al., ACRFence]** Y. Zheng, Y. Yang, W. Zhang, and A. Quinn, "ACRFence: Preventing semantic rollback attacks in agent checkpoint-restore," in *Proc. CoDAIM Workshop*, 2026. [Online]. Available: <https://arxiv.org/abs/2603.20625>

- **[Zheng et al., AgentSight]** Y. Zheng, Y. Hu, *et al.*, "AgentSight: System-level observability for AI agents using eBPF," in *Proc. PACMI @ SOSP*, 2025. [Online]. Available: <https://arxiv.org/abs/2508.02736>
- **[Zheng et al., CP-WBFT]** L. Zheng, J. Chen, Q. Yin, J. Zhang, X. Zeng, and Y. Tian, "Rethinking the reliability of multi-agent systems: A perspective from Byzantine fault tolerance," in *Proc. AAAI*, 2025. [Online]. Available: <https://arxiv.org/abs/2511.10400>
- **[Zheng et al., Judging LLM-as-a-Judge]** L. Zheng, *et al.*, "Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena," in *Proc. NeurIPS*, 2023. [Online]. Available: <https://arxiv.org/abs/2306.05685>
- **[Zhou et al., APE]** Y. Zhou, A. I. Muresanu, Z. Han, *et al.*, "Large language models are human-level prompt engineers (APE)," in *Proc. ICLR*, 2023. [Online]. Available: <https://arxiv.org/abs/2211.01910>
- **[Zhou et al., MASS]** H. Zhou, *et al.*, "Multi-agent design: Optimizing agents with better prompts and topologies (MASS)," in *Proc. ICLR*, 2026. [Online]. Available: <https://arxiv.org/abs/2502.02533>
- **[Zhou et al., SOTOPIA]** X. Zhou, H. Zhu, L. Mathur, *et al.*, "SOTOPIA: Interactive evaluation for social intelligence in language agents," in *Proc. ICLR Spotlight*, 2024. [Online]. Available: <https://arxiv.org/abs/2310.11667>
- **[Zhou et al., WebArena]** S. Zhou, F. F. Xu, H. Zhu, *et al.*, "WebArena: A realistic web environment for building autonomous agents," in *Proc. ICLR*, 2024. [Online]. Available: <https://arxiv.org/abs/2307.13854>
- **[Zhu et al., AgentDebug]** K. Zhu, Z. Liu, B. Li, *et al.*, "Where LLM agents fail and how they can learn from failures," in *Proc. ICML Workshop on Failure Modes of Agentic AI*, 2026. [Online]. Available: <https://arxiv.org/abs/2509.25370>
- **[Zhu et al., JudgeLM]** L. Zhu, X. Wang, and X. Wang, "JudgeLM: Fine-tuned large language models are scalable judges," in *Proc. ICLR*, 2025. [Online]. Available: <https://arxiv.org/abs/2310.17631>
- **[Zhu et al., MultiAgentBench]** K. Zhu, H. Du, Z. Hong, *et al.*, "MultiAgentBench: Evaluating the collaboration and competition of LLM agents," in *Proc. ACL*, 2025. [Online]. Available: <https://aclanthology.org/2025.acl-long.421/>
- **[Zhu et al., Orchestration Survey]** Y. Zhu, L. Liu, J. Yu, and D. Zhang, "LLM-based multi-agent orchestration: A survey of frameworks, communication protocols, and

emerging patterns," *Future Internet*, vol. 18, no. 6, art. 326, Jun. 2026. [Online]. Available: <https://doi.org/10.3390/fi18060326>

- **[Zhu et al., Rigorous Benchmarks]** Y. Zhu, T. Jin, *et al.*, "Establishing best practices in building rigorous agentic benchmarks," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2507.02825>
- **[Zhu et al., TraceLab]** K. Zhu, M. Jacob, C. Ma, Y. Pan, S. Wang, A. Krishnamurthy, and B. Kasikci, "TraceLab: Characterizing coding agent workloads for LLM serving," *arXiv: 2606.30560*, Jun. 2026. [Online]. Available: <https://arxiv.org/abs/2606.30560>
- **[Zhuge et al., Agent-as-a-Judge]** M. Zhuge, *et al.*, "Agent-as-a-Judge: Evaluate agents with agents," in *Proc. ICML*, 2025. [Online]. Available: <https://arxiv.org/abs/2410.10934>
- **[Zhuge et al., GPTSwarm]** M. Zhuge, *et al.*, "GPTSwarm: Language agents as optimizable graphs," in *Proc. ICML Oral*, 2024. [Online]. Available: <https://arxiv.org/abs/2402.16823>
- **[Zweiger et al., SEAL]** A. Zweiger, J. Pari, H. Guo, *et al.*, "Self-adapting language models (SEAL)," in *Proc. NeurIPS*, 2025. [Online]. Available: <https://arxiv.org/abs/2506.10943>